

breakpoint can be cleared by a breakpoint clear command, *c*, which can be preceded by labels and numbers, like the command *b*. There is a run command, *r*, in which the tracer executes until either a breakpoint, an exit call, or the end of the commands is encountered.

The tracer also keeps track of the subroutine level at which the program is running. This is shown in the window below the processor window and can also be seen through the indication numbers in the stack window. There are three commands that are based on these levels. The *-* command causes the tracer to run until the subroutine level is one less than the current level. What this command does is execute instructions until the current subroutine is finished. The converse is the *+* command, which runs the tracer until the next subroutine level is encountered. The *=* command runs until the same level is encountered, and can be used to execute a subroutine at the *CALL* command. If *=* is used, the details of the subroutine are not shown in the tracer window. There is a related command, *n*, which runs until the next line in the program is encountered. This command is especially useful when issued as a *LOOP* command; execution stops exactly when the bottom of the loop is executed.

## C.7 GETTING STARTED

In this section, we will explain how to use the tools. First of all, it is necessary to locate the software for your platform. We have precompiled versions for Solaris, UNIX, for Linux and for Windows. The tools are located on the CD-ROM and on the Web at [www.prenhall.com/tanenbaum](http://www.prenhall.com/tanenbaum). Once there, click on the *Companion Web Site* for this book and then click on the link in the left-hand menu. Unpack the selected zip file to a directory *assembler*. This directory and its subdirectories contain all the necessary material. On the CD-ROM, the main directories are *BigendNx*, *LtlendNx*, and *MSWindos*, and in each there is a subdirectory *assembler* which contains the material. The three top-level directories are for Big-Endian UNIX (e.g. Sun workstations), Little-Endian UNIX (e.g., Linux on PCs), and Windows systems, respectively.

After unpacking or copying, the assembler directory should contain the following subdirectories and files: *READ\_ME*, *bin*, *as\_src*, *trce\_src*, *examples*, and *exercise*. The precompiled sources can be found in the *bin* directory but, for convenience, there is also a copy of the binaries in the *examples* directory.

To get a quick preview of how the system works, go to the *examples* directory and type the command

```
t88 HelloWrld
```

This command corresponds to the first example in Sec. C.8.

The source code for the assembler is in the directory *as\_src*. The source code files are in the language C, and the command *make* should recompile the sources.

For POSIX-compliant platforms, there is a *Makefile* in the source directory which does the job. For Windows, there is a batch file *make.bat*. It may be necessary to move the executable files after compilation to a program directory, or to change the PATH variable to make the assembler *as88* and the tracer *t88* visible from the directories containing the assembly source codes. Alternatively, instead of typing *t88*, the full path name can be used.

On Windows 2000 and XP systems, it is necessary to install the *ansi.sys* terminal driver by adding the line

```
device=%systemRoot%\System32\ansi.sys
```

to the configuration file, *config.nt*. The location of this file is as follows:

Windows 2000: \winnt\system32\config.nt

Windows XP: \windows\system32\config.nt

On UNIX and Linux systems, the driver is usually standard.

## C.8 EXAMPLES

In Sec. C.2 through Sec. C.4, we discussed the 8088 processor, its memory, and its instructions. Then, in Sec. C.5, we studied the *as88* assembly language used in this tutorial. In Sec. C.6 we studied the tracer. Finally, in Sec. C.7, we described how to set up the toolkit. In theory, this information is sufficient to write and debug assembly programs with the tools provided. Nevertheless, it may be helpful for many readers to see some detailed examples of assembly programs and how they can be debugged with the tracer. That is the purpose of this section. All the example programs discussed in this section are available in the *examples* directory in the toolkit. The reader is encouraged to assemble and trace each one as it is discussed.

### C.8.1 Hello World Example

Let us start with the example of Fig. C-12, *HlloWrld.s*. The program is listed in the left window. Since the assembler's comment symbol is the exclamation mark (!), it is used in the program window to separate the instructions from the line numbers that follow. The first three lines contain constant definitions, which connect the conventional names of two system calls and the output file to their corresponding internal representations.

The pseudoinstruction .SECT, on line 4, states that the following lines should be considered to be part of the TEXT section; that is, processor instructions. Similarly, line 17 indicates that what follows is to be considered data. Line 19 initializes a string of data consisting of 12 bytes, including one space and a line feed (\n) at the end.

Lines 5, 18 and 20 contain labels, which are indicated by a colon :. These labels represent numerical values, similar to constants. In this case, however, the assembler has to determine the numerical values. Since *start* is at the beginning of the TEXT section, its value will be 0, but the value of any subsequent labels in the TEXT section (not present in this example), would depend on how many bytes of code preceded them. Now consider line 6. This line ends with the difference of two labels, which is numerically a constant. Thus, line 6 is effectively the same as

```
MOV CX,12
```

except that it lets the assembler determine the string length, rather than making the programmer do it. The value indicated here is the amount of space in the data reserved for the string on line 19. The MOV on line 6 is the copy command, which requires the *de - hw* to be copied to CX.

<pre>_EXI T = 1          ! 1 _WRITE = 4          ! 2 _STDOUT = 1         ! 3 .SECT .TEXT         ! 4 start:             ! 5     MOV CX,de-hw   ! 6     PUSH CX        ! 7     PUSH hw         ! 8     PUSH _STDOUT    ! 9     PUSH _WRITE     !10     SYS             !11     ADD SP,8        !12     SUB CX,AX       !13     PUSH CX        !14     PUSH _EXIT      !15     SYS             !16 .SECT .DATA         !17 hw:                !18 .ASCII "Hello World\n" !19 de:.BYTE 0          !20</pre>	<pre>CS: 00 DS=SS=ES: 002 AH:00 AL:0c AX: 12 BH:00 BL:00 BX: 0 CH:00 CL:0c CX: 12 DH:00 DL:00 DX: 0 SP: 7fd8 SF O D S Z C =&gt;0004 BP: 0000 CC - &gt; p - - 0001 =&gt; SI: 0000 IP:000c:PC 0000 DI: 0000 start + 7 000c</pre>	<pre>MOV CX,de-hw ! 6 PUSH CX      ! 7 PUSH HW      ! 8 PUSH _STDOUT ! 9 PUSH _WRITE  !10 SYS          !11 ADD SP,8     !12 SUB CX,AX   !13 PUSH CX     !14 SUB CX,AX   !15 PUSH CX     !16 PUSH _EXIT  !17 SYS          !18 &gt; Hello World\n hw + 0 = 0000: 48 65 6c 6f 20 57 6f Hello World 25928</pre>
----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

(a)

(b)

Figure C-12. (a) HlloWrld.s. (b) The corresponding tracer window.

Lines 7 through 11 show how system calls are made in the toolkit. These five lines are the assembly code translation of the C language function call

```
write(1, hw, 12);
```

where the first parameter is the file descriptor for standard output (1), the second is the address of the string to be printed (*hw*), and the third is the length of the string (12). Lines 7 through 9 push these parameters onto the stack in reverse order, which is the C calling sequence and the one used by the tracer. Line 10 pushes the system call number for write (4) onto the stack, and line 11 makes the actual call. While this calling sequence closely mimics how an actual assembly language program would work on a UNIX (or Linux) PC, for a different operating system, it would have to be modified slightly to use the calling conventions of that operating

system. The *as88* assembler and *t88* tracer use the UNIX calling conventions even when they are running on Windows, however.

The system call on line 11 does the actual printing. Line 12 performs a cleanup on the stack, resetting the stack pointer back to the value it had before the four 2-byte words were pushed onto the stack. If the write call is successful, the number of bytes written is returned in AX. Line 13 subtracts the system call result after line 11 from the original string length in CX to see whether the call was successful, that is, to see if all the bytes were written. Thus, the exit status of the program will be 0 on success and something else on failure. Lines 14 and 15 prepare for the `exit` system call on line 16 by pushing the exit status and function code for the `EXIT` call onto the stack.

Note that in the `MOV` and `SUB` instructions the first argument is the destination and the second is the source. This is the convention used by our assembler; other assemblers may reverse the order. There is no particular reason to choose one order over the other.

Now let us try to assemble and run *HlloWrld.s*. Instructions will be given for both UNIX and Windows platforms. For Linux, Solaris, Mac OS X, and other UNIX variants, the procedure should be essentially the same as for UNIX. First, start up a command prompt (shell) window. On Windows, the click sequence is usually

Start > Programs > Accessories > Command prompt

Next, change to the *examples* directory using the `cd` (Change directory) command. The argument to this command depends on where the toolkit has been placed in the file system. Then verify that the assembler and tracer binaries are in this directory, using `ls` on UNIX and `dir` on Windows systems. They are called *as88* and *t88*, respectively. On Windows systems, they have the extension `.exe`, but that need not be typed in the commands. If the assembler and tracer are not there, find them and copy them there.

Now assemble the test program using

`as88 HlloWrld.s`

If the assembler is present in the *examples* directory but this command gives an error message, try typing

`./as88 HlloWrld.s`

on UNIX systems or

`.\as88 HlloWrld.s`

on Windows systems.

If the assembly process completes correctly, the following messages will be displayed:

`Project HlloWrld listfile HlloWrld.$`

```
Project HlloWrld num file HlloWrld.#  
Project HlloWrld loadfile HlloWrld.88.
```

and the corresponding three files created. If there are no error messages, give the tracer command:

```
t88 HlloWrld
```

The tracer display will appear with the arrow in the upper right-hand window pointing to the

```
MOV CX,de-hw
```

instruction of line 6. Now hit the return (called Enter on PC keyboards) key. Notice that the instruction pointed to is now

```
PUSH CX
```

and the value of CX in the left-hand window is now 12. Hit return again and notice that the middle window on the top line now contains the value 000c, which is hex-adecimal for 12. This window shows the stack, which now has one word containing 12. Now hit return three more times to see the PUSH instructions on lines 8, 9, and 10 being carried out. At this point, the stack will have four items and the program counter in the left-hand window will have the value 000b.

The next time return is hit, the system call is executed and the string “Hello World\n” is displayed in the lower right-hand window. Note that SP now has the value 0x7ff0. After the next return, SP is incremented by 8 and becomes 0x7ff8. After four more returns, the exit system call completes and the tracer exits.

To be certain that you understand how everything works, fetch the file *hllowrld.s* into your favorite editor. It is better not to use a word processor. On UNIX systems, *ex*, *vi*, or *emacs* are good choices. On Windows systems, *notepad* is a simple editor, usually reachable from

```
Start > Programs > Accessories > Notepad
```

Do not use *Word* since the display will not look right and the output may be formatted incorrectly.

Modify the string on line 19 to display a different message, then save the file, assemble it, and run it with the tracer. You are now starting to do assembly language programming.

### C.8.2 General Registers Example

The next example demonstrates in more detail how the registers are displayed and one of the pitfalls of multiplication on the 8088. In Fig. C-13, part of the program *genReg.s* is shown on the left. To its right are two tracer register windows, corresponding to different stages of the program’s execution. Fig. C-13(b) shows the register state after line 7 has been executed. The instruction

```
MOV AX,258
```

on line 4 loads the value 258 in AX, which results in the value 1 being loaded into

AH and the value 2 being loaded into AL. Then line 5 adds AL to AH, making AH equal to 3. On line 6, the contents of the variable *times* (10) are copied into CX. On line 7, the address of the variable *muldat*, which is 2 because it is at the second byte of the DATA segment, is loaded into BX. This is the instant in time at which the dump of Fig. C-13(b) was made. Note that AH is 3, AL is 2, and AX is 770, which is to be expected, as  $3 \times 256 + 2 = 770$ .

<pre> start:           ! 3     MOV  AX,258   ! 4     ADDB AH,AL    ! 5     MOV  CX,(times) ! 6     MOV  BX,muldat ! 7     MOV  AX,(BX)   ! 8     l1p: MUL  2(BX) ! 9     LOOP l1p       !10     .SECT .DATA      !11     times: .WORD 10   !12     muldat:.WORD 625,2 !13 </pre>	<pre> CS: 00 DS=SS=ES002 AH:03 AL:02 AX: 770 BH:00 BL:02 BX:  2 CH:00 CL:0a CX: 10 DH:00 DL:00 DX:  0 SP: 7fe0 SF O D S Z C BP: 0000 CC - &gt; p - - SI: 0000 IP:0009:PC DI: 0000 start + 4 </pre>	<pre> CS: 00 DS=SS=ES002 AH:38 AL:80 AX: 14464 BH:00 BL:02 BX:  2 CH:00 CL:04 CX:  4 DH:00 DL:01 DX:  1 SP: 7fe0 SF O D S Z C BP: 0000 CC v &gt; p - c SI: 0000 IP:0011:PC DI: 0000 start + 7 </pre>
(a)	(b)	(c)

**Figure C-13.** (a) Part of a program. (b) The tracer register window after line 7 has been executed. (c) The registers.

The next instruction (line 8) copies the contents of *muldat* into AX. Thus, after the return key is hit, AX will be 625.

We are now ready to enter a loop that multiplies the contents of AX by the word addressed by 2BX (i.e., *muldat* + 2), which has the value 2. The implied destination of the MUL instruction is the DX : AX long register combination. In the first iteration of the loop, the result fits in one word, so AX contains the result (1250), and DX remains 0. The contents of all the registers after 7 multiplications are shown in Fig. C-13.

Since AX started at 625, the result after those seven multiplications by 2 is 80,000. This result does not fit in AX, but the product is held in the 32-bit register formed by the concatenation of DX : AX, so DX is 1 and AX is 14,464. Numerically, this value is  $1 \times 65,536 + 14,464$ , which is, indeed, 80,000. Note that CX is 4 here, because the LOOP instruction decrements it every iteration. Because it started at 10, after seven executions of the MUL instruction (but only six iterations of the LOOP instruction) we have CX set to 4.

In the next multiplication, trouble crops up. Multiplication involves AX but not DX, so the MUL multiples AX (14464) by 2 to get 28,928. This results in AX being set to 28,928 and DX being set to 0, which is numerically incorrect.

### C.8.3 Call Command and Pointer Registers

The next example, *vecprod.s* is a small program that computes the inner product of two vectors, *vec1* and *vec2*. It is listed in Fig. C-14.

The first part of the program prepares to call *vecmul* by saving SP in BP and then pushing the addresses of *vec2* and *vec1* onto the stack so that *vecmul* will have access to them. Then the length of the vector in bytes is loaded in CX on line 8. By shifting this result right one bit, on line 9, CX now contains the number of words in the vector, which is pushed onto the stack on line 10. The call to *vecmul* is made on line 11.

Once again, it is worth mentioning that the arguments of subroutines are, by convention, pushed onto the stack in reverse order to be compatible with the C calling convention. In this way, *vecmul* can also be called from C using

```
vecmul(count, vec1, vec2)
```

During the CALL instruction, the return address is pushed onto the stack. If the program is traced, then this address turns out to be 0x0011.

The first instruction in the subroutine is a PUSH of the base pointer, BP, on line 22. BP is saved because we will need this register to address the arguments and the local variables of the subroutine. Next, the stack pointer is copied to the BP register on line 23, so that the new value of the base pointer is pointing to the old value.

Now everything is ready for loading the arguments into registers and for reserving space for a local variable. In the next three lines, each of the arguments is fetched from the stack and put in a register. Recall that the stack is word oriented, so stack addresses should be even. The return address is next to the old base pointer so it is addressed by 2(BP). The *count* argument is next and addressed by 4(BP). It is loaded into CX on line 24. In lines 25 and 26, SI is loaded with *vec1* and DI is loaded with *vec2*. This subroutine needs one local variable with initial value 0 to save the intermediate result, so the value 0 is pushed on line 27.

The state of the processor just before the loop is entered on line 28 for the first time is shown in Fig. C-15. The narrow window in the middle of the top row (to the right of the registers) shows the stack. At the bottom of the stack is the address of *vec2* (0x0022), with *vec1* (0x0018) above it and the third argument, the number of items in each vector (0x0005) above that. Next comes the return address (0x0011). The number 1 to the left of this address indicates it is a return address one level from the main program. In the window below the registers, the same number 1 is shown, this time giving its symbolic address. Above the return address in the stack is the old value of BP (0x7fc0) and then the zero pushed on line 27. The arrow pointing to this value indicates where SP points. The window to the right of the stack shows a fragment of the program text, with the arrow indicating the next instruction to be executed.

Now let us examine the loop starting at line 28. The instruction LODS loads a memory word indirectly through the register SI from the data segment into AX. Because the direction flag is set, LODS is in auto-increment mode, so after the instruction SI will point to the next entry of *vec1*.

To see this effect graphically, start the tracer with the command

```
t88 vecprod
```

```

_EXIT = 1           ! 1 define the value of _EXIT
_PRINTF = 127      ! 2 define the value of _PRINTF
.SECT .TEXT
instart:
    MOV BP,SP
    PUSH vec2
    PUSH vec1
    MOV CX,vec2-vec1
    SHR CX,1
    PUSH CX
    CALL vecmul
    MOV (inprod),AX
    PUSH AX
    PUSH pfmt
    PUSH _PRINTF
    SYS
    ADD SP,12
    PUSH 0
    PUSH _EXIT
    SYS

    vecmul:
        PUSH BP
        MOV BP,SP
        MOV CX,4(BP)
        MOV SI,6(BP)
        MOV DI,8(BP)
        PUSH 0
        LODS
        MUL (DI)
        ADD -2(BP),AX
        ADD DI,2
        LOOP 1b
        POP AX
        POP BP
        RET

.SECT .DATA
pfmt: .ASCIZ "Inner product is: %d\n"
.ALIGN 2
vec1:.WORD 3,4,7,11,3
vec2:.WORD 2,6,3,1,0
.SECT .BSS
inprod: .SPACE 2

```

! 1 define the value of \_EXIT  
! 2 define the value of \_PRINTF  
! 3 start the TEXT segment  
! 4 define label instart  
! 5 save SP in BP  
! 6 push address of vec2  
! 7 push address of vec1  
! 8 CX = number of bytes in vector  
! 9 CX = number of words in vector  
! 10 push word count  
! 11 call vecmul  
! 12 move AX  
! 13 push result to be printed  
! 14 push address of format string  
! 15 push function code for PRINTF  
! 16 call the PRINTF function  
! 17 clean up the stack  
! 18 push status code  
! 19 push function code for EXIT  
! 20 call the EXIT function  
  
! 21 start of vecmul(count, vec1, vec2)  
! 22 save BP on stack  
! 23 copy SP into BP to access arguments  
! 24 put count in CX to control loop  
! 25 SI = vec1  
! 26 DI = vec2  
! 27 push 0 onto stack  
  
! 28 move (SI) to AX  
! 29 multiply AX by (DI)  
! 30 add AX to accumulated value in memory  
! 31 increment DI to point to next element  
! 32 if CX > 0, go back to label 1b  
! 33 pop top of stack to AX  
! 34 restore BP  
! 35 return from subroutine  
  
! 36 start DATA segment  
! 37 define string  
! 38 force address even  
! 39 vector 1  
! 40 vector 2  
! 41 start BSS segment  
! 42 allocate space for inprod

Figure C-14. The program *vecprod.s*.

When the tracer window appears, type the command

```
/vecmul+7b
```

followed by a return to put a breakpoint at the line containing the LODS. From now on, we will not mention that all commands must be followed by the return key. Then give the command

```
g
```

to have the tracer execute commands until the breakpoint is encountered. It will stop at the line containing the LODS.

On line 29, the value of AX is multiplied to the source operand. The memory word for the MUL instruction is fetched from the data segment through the DI in register indirect mode. The implied destination of MUL is the DX : AX long register combination which is not mentioned in the instruction but which is implied by it.

On line 30, the result is added to the local variable at the stack address -2(BP). Because MUL does not autoincrement its operand, that must be done explicitly on line 31. Afterward, DI points to the next entry of *vec2*.

The LOOP instruction finishes this step. Register CX is decremented, and, if it is still positive, the program jumps back to the local label *l* on line 28. The use of the local label *lb* means the closest label *l* looking backward from the current location. After the loop, the subroutine pops the return value into AX (line 33), restores BP (line 34), and returns to the calling program (line 35).

Then the main program is resumed after the call with the MOV instruction on line 12. This instruction is the start of a five-instruction sequence whose goal is to print the result. The printf system call is modeled after the *printf* function in the standard C programming library. Three arguments are pushed onto the stack on lines 13-15. These arguments are the integer value to be printed, the address of the format string (*pfmt*), and the function code for printf (127). Note that the format string *pfmt* contains a %d to indicate that an integer variable can be found as argument to the printf call to complete the output.

Line 17 cleans up the stack. Since the program started on line 5 by saving the stack pointer in the base pointer, we could also use the instruction

```
MOV SP,BP
```

for a stack cleanup. The advantage of this solution is that the programmer does not need to keep the stack balanced in the process. For the main program this is not a big issue, but in subroutines this approach is an easy way to throw away garbage such as obsolete local variables.

The subroutine *vecmul* can be included in other programs. If the source file *vecprod.s* is put on the command line behind another assembler source file, the subroutine is available for multiplying two vectors of a fixed length. It is advisable to remove the constant definitions \_EXIT and \_PRINTF first, in order to avoid their being defined twice. If the header file *syscalnr.h* is included somewhere, then there is no need to define the system call constants anywhere else.

MOV BP,SP	! 5	CS: 00 DS=SS=ES004		PUSH BP	! 22
PUSH vec2	! 6	AH:00 AL:00 AX: 0		MOV BP,SP	! 23
PUSH vec1	! 7	BH:00 BL:00 BX: 0		MOV CX,4(BP)	! 24
MOV CX,vec2-vec1	! 8	CH:00 CL:05 CX: 5 =>0000		MOV SI,6(BP)	! 25
SHR CX,1	! 9	DH:00 DL:00 DX: 0 7fc0		MOV DI,8(BP)	! 26
PUSH CX	! 10	SP: 7fb4 SF O D S Z C 1 0011		PUSH 0	! 27
CALL vecmul	! 11	BP: 7fb6 CC - > p z - 0005 =>1:		LODS	! 28
-----		SI: 0018 IP:0031:PC 0018		MUL (DI)	! 29
vecmul :	! 21	DI: 0022 vecmul+7 0022		ADD -2(BP),AX	! 30
PUSH BP	! 22				
MOV BP,SP	! 23				
MOV CX,4(BP)	! 24				
MOV SI,6(BP)	! 25				
MOV DI,8(BP)	! 26				
PUSH 0	! 27				
1: LODS	! 28				
MUL (DI)	! 29				
ADD -2(BP),AX	! 30				
ADD DI,2	! 31				
LOOP 1b	! 32				

1 <= inpstart + 7  
■ >

```

vec1+0 = 0018: 3 0 4 0 7 0 b 0 ..... 3
vec2+0 = 0022: 2 0 6 0 3 0 1 0 ..... 2
pfmt+0 = 0000:54 68 65 20 69 6e 20 70 The in prod 26708
pfmt+18 = 0012:25 64 21 a 0 0 3 0 % d!.....25637

```

**Figure C-15.** Execution of *vecprod.s* when it reaches line 28 for the first time.

### C.8.4 Debugging an Array Print Program

In the previous examples, the programs examined were simple but correct. Now we will show how the tracer can help debug incorrect programs. The next program is supposed to print the integer array, which is supplied after the label *vec1*. However, the initial version contains three errors. The assembler and tracer will be used to correct those errors, but first we will discuss the code.

Because every program needs system calls, and thus must define constants by which to identify the call numbers, we have put the constant definitions for those numbers in a separate header file *./syscalnr.h*, which is included on line 1 of the code. This file also defines the constants for the file descriptors

```

STDIN = 0
STDOUT = 1
STDERR = 2

```

which are opened at the start of the process, and header labels for the text and the data segments. It is sensible to include it at the head of all assembly source files, as these are much used definitions. If a source is distributed over more than one file, the assembler includes only the first copy of this header file, to avoid defining the constants more than once.

The program *arrayprt* is shown in Fig. C-16. Comments have been omitted here, as the instructions should be well known by now. This layout allows a two-column format. Line 4 puts the address of the empty stack in the base pointer register to allow the stack cleanup can be made on line 10 by copying the base pointer to the stack pointer, as described in the previous example. We also have seen the computation and pushing of the stack arguments before the call on lines 5 through 9 in the previous example. Lines 22 to 25 load the registers in the subroutine.

```

#include "../syscalnr.h"      ! 1           .SECT .TEXT          ! 20
                               ! 2           vecprint:          ! 21
.SECT .TEXT                 ! 3           PUSH BP            ! 22
vecpstr:                   ! 4           MOV  BP,SP          ! 23
    MOV  BP,SP             ! 5           MOV  CX,4(BP)       ! 24
    PUSH vec1              ! 6           MOV  BX,6(BP)       ! 25
    MOV  CX,frmatstr-vec1 ! 7           MOV  SI,0            ! 26
    SHR  CX                ! 8           PUSH frmatkop      ! 27
    PUSH CX                ! 9           PUSH frmatstr      ! 28
    CALL vecprint          ! 10          PUSH _PRINTF        ! 29
    MOV  SP,BP              ! 11          SYS               ! 30
    PUSH 0                 ! 12          MOV  -4(BP),frmatint ! 31
    PUSH _EXIT              ! 13          1: MOV  DI,(BX)(SI)   ! 32
    SYS                     ! 14          INC   SI            ! 35
                               ! 15          LOOP 1b            ! 36
vec1: .WORD 3,4,7,11,3     ! 15          PUSH '\n'          ! 37
frmatstr: .ASCII "%s"      ! 16          PUSH _PUTCHAR       ! 38
                               ! 17          SYS               ! 39
frmatkop:                  ! 17          MOV  SP,BP          ! 40
.ASCIZ "The array contains " ! 18          RET               ! 41
frmatint: .ASCII "%d"      ! 19

```

**Figure C-16.** The program *arrayprt* before debugging.

Lines 27 to 30 show how a string can be printed, and 31 to 34 show the printf system call for an integer value. Note that the address of the string is pushed on line 27, while on line 33 the value of the integer is moved onto the stack. In both cases the address of the format string is the first argument of PRINTF. Lines 37 to 39 show how a single character can be printed using the putchar system call.

Now let us try assembling and running the program. When the command

```
as88 arrayprt.s
```

is typed, we get an operand error on line 28 of the file *arrayprt.\$*. This file is generated by the assembler by combining the included files with the source file to get a composite file that is the actual assembler input. To see where line 28 really is, we have to examine line 28 of *arrayprt.\$*. We cannot look at *arrayprt.s* to get the line number because the two files do not match on account of the header being included line by line in *arrayprt.\$*. Line 28 in *arrayprt.\$* corresponds to line 7 in *arrayprt.s* because the included header file, *syscalnr.h*, contains 21 lines.

One easy way to find line 28 of *arrayprt.\$* on UNIX is to type the command

```
head -28 arrayprt.$
```

which displays the first 28 lines of the combined file. The line at the bottom of the listing is the one in error. In this way (or by using an editor and going to line 28)

we see that the error is on line 7, which contains the `SHR` instruction. Comparing this code with the instruction table in Fig. C-4 shows the problem: the shift count has been omitted. The corrected line 7 should read

```
SHR CX,1
```

It is very important to note that the error must be corrected in the original source file, *arrayprt.s*, and *not* in the combined source, *arrayprt.\$*, as the latter is automatically regenerated every time the assembler is called.

The next attempt to assemble the source code file should succeed. Then the tracer can be started by the command:

```
t88 arrayprt
```

During the tracing process, we can see that the output is not consistent with the vector in the data segment. The vector contains: 3, 4, 7, 11, 3, but the values displayed start with: 3, 1024, . . . Clearly, something is wrong.

To find the error, the tracer can be run again, step by step, examining the state of the machine just before the incorrect value is printed. The value to be printed is stored in memory on lines 32 and 33. Since the wrong value is being printed, this is a good place to see what is wrong. The second time through the loop, we see that `SI` is an odd number, when clearly it should be an even number, as it is indexing through words, not bytes. The problem is on line 35. It increments `SI` by 1; it should increment it by 2. To fix the bug, this line should be changed to

```
ADD SI,2
```

When this correction is made, the printed list of numbers is correct.

However, there is one more error waiting for us. When *vecprint* is finished and returns, the tracer complains about the stack pointer. The obvious thing to check for now is whether the value pushed onto the stack when *vecprint* is called is the value on top of the stack when the `RET` on line 41 is executed. It is not. The solution is to replace line 40 with two lines:

```
ADD SP,10  
POP BP
```

The first instruction removes the 5 words pushed onto the stack during *vecprint*, thus exposing the value of `BP` saved on line 22. By popping this value to `BP`, we restore `BP` to its precall value and expose the correct return address. Now the program terminates correctly. Debugging assembly code is definitely more of an art than a science, but the tracer makes it much easier than running on the bare metal.

### C.8.5 String Manipulation and String Instructions

The main purpose of this section is to show how to handle repeatable string instructions. In Fig. C-17, there are two simple string manipulation programs, *strngcpy.s* and *reversprs.s*, both present in the *examples* directory. The one in

.SECT .TEXT		#include "../syscalnr.h"	
stcstart:	! 1		! 1
PUSH mesg1	! 2	start: MOV DI,str	! 2
PUSH mesg2	! 3	PUSH AX	! 3
CALL strngcpy	! 4	MOV BP,SP	! 4
ADD SP,4	! 5	PUSH _PUTCHAR	! 5
PUSH 0	! 6	MOVB AL,'n'	! 6
PUSH 1	! 7	MOV CX,-1	! 7
SYS	! 8	REP NZ SCASB	! 8
strngcpy:	! 9	NEG CX	! 9
PUSH CX	! 10	STD	! 10
PUSH SI	! 11	DEC CX	! 11
PUSH DI	! 12	SUB DI,2	! 12
PUSH BP	! 13	MOV SI,DI	! 13
MOV BP,SP	! 14	1: LODSB	! 14
MOV AX,0	! 15	MOV (BP),AX	! 15
MOV DI,10(BP)	! 16	SYS	! 16
MOV CX,-1	! 17	LOOP 1b	! 17
REPNZ SCASB	! 18	MOVB (BP),'n'	! 18
NEG CX	! 19	SYS	! 19
DEC CX	! 20	PUSH 0	! 20
MOV SI,10(BP)	! 21	PUSH _EXIT	! 21
MOV DI,12(BP)	! 22	SYS	! 22
PUSH DI	! 23	.SECT .DATA	! 23
REP MOVS B	! 24	str: .ASCIZ "reverse\n"	! 24
CALL stringpr	! 25		
MOV SP,BP	! 26		
POP BP	! 27		
POP DI	! 28		
POP SI	! 29		
POP CX	! 30		
RET	! 31		
.SECT .DATA	! 32		
mesg1: .ASCIZ "Have a look\n"	! 33		
mesg2: .ASCIZ "qrst\n"	! 34		
.SECT .BSS			

(a)

(b)

**Figure C-17.** (a) Copy a string (*strngcpy.s*). (b) Print a string backward (*reverspr.s*).

Fig. C-17(a) is a subroutine for copying a string. It calls a subroutine, *stringpr*, which can also be found in a separate file *stringprs.s*. It is not listed in this appendix. In order to assemble programs containing subroutines in separate source files, just list all source files in the *as88* command, starting with the source file for

the main program, which determines the names of the executable and the auxiliary files. For example, for the program of Fig. C-17(a) use

```
as88 strngcpy.s stringpr.s
```

The program of Fig. C-17(b) outputs strings in reverse order. We will now look at them in turn.

To demonstrate that the line numbers are really just comments, in Fig. C-17(a) we have numbered the lines starting with the first label, omitting what comes before them. The main program, on lines 2 through 8, first calls *strngcpy* with two arguments, the source string, *mesg2*, and the destination string, *mesg1*, in order to copy the source to the destination.

Now let us look at *strngcpy*, starting on line 9. It expects that the addresses of the destination buffer and the string source have been pushed onto the stack just before the subroutine is called. On lines 10 to 13, the registers used are saved by pushing them onto the stack so that they can be restored later on lines 27 to 30. On line 14, we copy SP to BP in the usual way. Now BP can be used to load the arguments. Again, on line 26, we clean the stack by copying BP to SP.

The heart of the subroutine is the instruction REP MOVSB, on line 24. The instruction MOVSB moves the byte pointed to by SI to the memory address pointed to by DI. Both SI and DI are then incremented by 1. The REP creates a loop in which this instruction is repeated, decrementing CX by 1 for each byte moved. The loop is terminated when CX reaches 0.

Before we can run the REP MOVSB loop, however, we have to set up the registers, which is done in lines 15 through 22. The source index, SI, is copied from the argument on the stack on line 21; the destination index, DI, is set up on line 22. Obtaining the value of CX is more involved. Note that the end of a string is indicated by a zero byte. The MOVSB instruction does not affect the zero flag, but the instruction SCASB (scan byte string) does. It compares the value pointed to by DI with the value in AL, and it increments DI on the fly. Moreover, it is repeatable like MOVSB. So, on line 15 AX and hence AL is cleared, on line 16 the pointer for DI is fetched from the stack, and CX is initialized to -1 on line 17. On line 18, we have the REPNZ SCASB, which does the comparison in loop context, and sets the zero flag on equality. At each step of the loop, CX is decremented, and the loop stops when the zero flag is set, because the REPNZ checks both the zero flag and CX. The number of steps for the MOVSB loop is now computed as the difference of the current value of CX and the previous -1 on lines 19 and 20.

It is cumbersome that there are two repeatable instructions necessary, but this is the price for the design choice that move instructions never affect the condition codes. During the loops, the index registers have to be incremented, and to this end it is necessary that the direction flag is cleared.

Lines 23 and 25 print the copied string by means of a subroutine, *stringpr*, which is in the *examples* directory. It is straightforward and will not be discussed here.

In the reverse print program shown in Fig. C-17(b), the first line includes the usual system call numbers. On line 3, a dummy value is pushed onto the stack, and on line 4, the base pointer, BP, is made to point to the current top of stack. The program is going to print ASCII characters one by one, thus the numerical value \_PUTCHAR is pushed onto the stack. Note that BP points to the character to be printed when a SYS call is made.

Line 2, 6 and 7 prepare the registers DI, AL and CX for the repeatable SCASB instruction. The count register and the destination index are loaded in a similar way as in the string copy routine, but the value of AL is the new line character, instead of the value 0. In this way, the SCASB instruction will compare the character values of the string *str* to \n instead of to 0, and set the zero flag whenever it is found.

The REP SCASB increments the DI register, so, after a hit, the destination index points at the zero character following the new line. On line 12, DI is decremented by two to have it point to the last letter of the word.

If the string is scanned in reverse order and printed character by character, we have obtained our goal, so on line 10 the direction flag is set to reverse the adjustment of the index registers in the string instructions. Now the LODSB on line 14 copies the character in AL, and on line 15 this character is put just next to the \_PUTCHAR on the stack, so the SYS instruction prints it.

The instructions on lines 18 and 19 print an additional new line and the program closes with an \_EXIT call in the usual way.

The current version of the program contains a bug. It can be found if the program is traced step by step.

The command /str will put the string *str* in the tracer data field. Since the numerical value of the data address is also given, we can find out how the index registers run through the data with respect to the position of the string.

The bug, however, is encountered only after hitting the return many times. By using the tracer commands we can get to the problem faster. Start the tracer and give the command 13 to put us in the middle of the loop. If we now give the command b we set a breakpoint on this line 15. If we give two new lines, then we see that the final letter e is printed in the output field. The r command will keep the tracer running until either a breakpoint or the end of the process is encountered. In this way, we can run through the letters by giving the r command repeatedly until we are close to the problem. From this point, we can run the tracer at one step at a time until we see what happens at the critical instructions.

We can also put the breakpoint at a specific line, but then we must keep in mind, that the file ..//syscalnr is included, which causes the line numbers to be offset by 20. Consequently, the breakpoint on line 16 can be set by the command 36b. This is not an elegant solution, so it is much better to use the global label start on line 2 before the instruction and give the command /start+14b, which puts the breakpoint in the same place without having to keep track of the size of the included file.

### C.8.6 Dispatch Tables

In several programming languages, there exist *case* or *switch* statements to select a jump from several alternatives according to some numerical value of a variable. Sometimes, such multiway branches are also needed in assembly language programs, too. Think, for instance, of a set of system call subroutines combined in a single `SYS` trap routine. The program `jumptbl.s`, shown in Fig. C-18 shows how such a multi-branch switch can be programmed in 8088 assembler.

The program starts by printing the string whose label is `strt`, inviting the user to type an octal digit (lines 4 through 7). Then a character is read from standard input (lines 8 and 9). If the value in `AX` is less than 5, the program interprets it as an end of file marker and jumps to the label `8` on line 22 to exit with a status code of 0.

If end of file has not been encountered, the incoming character, in `AL`, is inspected. Any character less than the digit 0 is considered to be white space and is ignored by the jump on line 13, which retrieves another character. Any character over digit 9 is considered to be incorrect input. On line 16, it is mapped onto the ASCII colon character, which is the successor of digit 9 in the ASCII character sequence.

Thus, on line 17 we have a value in `AX` between digit 0 and the colon. This value is copied into `BX`. On line 18, the `AND` instruction masks off all but the low-order four bits, which leaves the number between 0 and 10 (due to the fact that ASCII 0 is `0x30`). Since we are going to index into a table of words, rather than bytes, the value in `BX` is multiplied by two using the left shift on line 19.

On line 20, we have a `call` instruction. The effective address is found by adding the value of `BX` to the numerical value of label `tbl`, and the contents of this composite address are loaded into the program counter, `PC`.

This program chooses one out of ten subroutines according to a character which is fetched from standard input. Each of those subroutines pushes the address of some message onto the stack and then jumps to a `_PRINTF` system subroutine call which is shared by all of them.

In order to understand what is happening, we need to be aware that the `JMP` and `CALL` instructions load some text segment address in `PC`. Such an address is just a binary number, and during the assembly process all addresses are replaced by their binary values. Those binary values can be used to initialize an array in the data segment, and this is done in line 50. Thus, the array starting at `tbl` contains the starting addresses of `rout0`, `rout1`, `rout2`, and so on, two bytes per address. The need for 2-byte addresses explains why we needed the 1-bit shift on line 19. A table of this type is often called a **dispatch table**.

How those routines work can be seen in the `erout` routine on lines 43 through 48. This routine handles the case of an out-of-range digit. First, the address of the message (in `AX`) is pushed onto the stack on line 43. Then the number of the `_PRINTF` system call is pushed onto the stack. After that, the system call is made, the stack is cleaned up, and the routine returns. The other nine routines, `rout0`

```

#include "../syscalnr.h" ! 1
.SECT .TEXT ! 2
jumpstrt: ! 3
    PUSH strt ! 4
    MOV BP,SP ! 5
    PUSH _PRINTF ! 6
    SYS ! 7
    PUSH _GETCHAR ! 8
1: SYS ! 9
    CMP AX,5 ! 10
    JL 8f ! 11
    CMPB AL,'0' ! 12
    JL 1b ! 13
    CMPB AL,'9' ! 14
    JLE 2f ! 15
    MOVB AL,'9'+1 ! 16
2: MOV BX,AX ! 17
    AND BX,0Xf ! 18
    SAL BX,1 ! 19
    CALL tbl(BX) ! 20
    JMP 1b ! 21
8: PUSH 0 ! 22
    PUSH _EXIT ! 23
    SYS ! 24
rout0: MOV AX,mes0 ! 25
                JMP 9f ! 26
rout1: MOV AX,mes1 ! 27
                JMP 9f ! 28
rout2: MOV AX,mes2 ! 29
                JMP 9f ! 30
rout3: MOV AX,mes3 ! 31
                JMP 9f ! 32
rout4: MOV AX,mes4 ! 33
                JMP 9f ! 34
rout5: MOV AX,mes5 ! 35
                JMP 9f ! 36
rout6: MOV AX,mes6 ! 37
                JMP 9f ! 38
rout7: MOV AX,mes7 ! 39
                JMP 9f ! 40
rout8: MOV AX,mes8 ! 41
                JMP 9f ! 42
erout: MOV AX,emes ! 43
9: PUSH AX ! 44
    PUSH _PRINTF ! 45
SYS ! 46
ADD SP,4 ! 47
RET ! 48

.SECT .DATA ! 49
tbl: WORD rout0,rout1,rout2,rout3,rout4,rout5,rout6,rout7,rout8,rout8,erout ! 50
mes0: .ASCIZ "This is a zero.\n" ! 51
mes1: .ASCIZ "How about a one.\n" ! 52
mes2: .ASCIZ "You asked for a two.\n" ! 53
mes3: .ASCIZ "The digit was a three.\n" ! 54
mes4: .ASCIZ "You typed a four.\n" ! 55
mes5: .ASCIZ "You preferred a five.\n" ! 56
mes6: .ASCIZ "A six was encountered.\n" ! 57
mes7: .ASCIZ "This is number seven.\n" ! 58
mes8: .ASCIZ "This digit is not accepted as an octal.\n" ! 59
emes: .ASCIZ "This is not a digit. Try again.\n" ! 60
strt: .ASCIZ "Type an octal digit with a return. Stop on end of file.\n" ! 61

```

**Figure C-18.** A program demonstrating a multiway branch using a dispatch table.

through *rout8*, each load the addresses of their private messages in AX, and then jump to the second line of *erout* to output the message and finish the subroutine.

In order to get accustomed to the dispatch tables, the program should be traced with several different input characters. As an exercise, the program can be changed

in such a way that all characters generate a sensible action. For example, all characters other than the octal digits should give an error message.

### C.8.7 Buffered and Random File Access

The program *InFilBuf.s*, shown in Fig. C-19, demonstrates random I/O on files. A file is assumed to consist of some number of lines, with different lines potentially having different lengths. The program first reads the file and builds a table in which entry *n* is the file position at which line *n* begins. Afterward, a line can be requested, its position looked up in the table, and the line read in by means of *Iseek* and *read* system calls. The file name is given as the first input line on standard input. This program contains several fairly independent chunks of code, which can be modified for other purposes.

The first five lines simply define the system call numbers and the buffer size, and set the base pointer at the top of the stack, as usual. Lines 6 through 13 read the file name from standard input, and store it as a string at label *linein*. If the file name is not properly closed with a new line, then an error message is generated, and the process exits with a nonzero status. This is done in lines 38 through 45. Note that the address of the file name is pushed on line 39, and the address of an error message is pushed on line 40. If we examine the error message itself, (on line 113) then we have a *%s* string request in the *\_PRINTF* format. The contents of the string *linein* are inserted here.

If the file name can be copied without problems, the file is opened on lines 14 to 20. If the *open* call fails, then the return value is negative and a jump is made to the label 9 on line 28 to print an error message. If the system call succeeds, then the return value is a file descriptor, which is stored in the variable *fildes*. This file descriptor is needed in the subsequent *read* and *Iseek* calls.

Next, we read the file in blocks of 512 bytes, each of which is stored in the buffer *buf*. The buffer allocated is two bytes larger than the necessary 512 bytes, just to demonstrate how a symbolic constant and a integer can be mixed in an expression (on line 123). In the same way, on line 21 *SI* is loaded with the address of the second element of the *linh* array, which leaves a machine word containing 0 at the bottom of this array. The register *BX* will contain the file address of the first unread character of the file, and hence, it is initialized to 0 before the first time that the buffer is filled on line 22.

The filling of the buffer is handled by the *fillbuf* routine on lines 83 through 93. After pushing the arguments for the *read*, the system call is requested, which puts the number of characters actually read in *AX*. This number is copied into *CX* and the number of characters still in the buffer will be kept in *CX* thereafter. The file position of the first unread character in the file is kept in *BX*, so *CX* has to be added to *BX* in line 91. On line 92, the buffer bottom is put into *DI* in order to get ready to scan the buffer for the next new line character.

After returning from *fillbuf*, line 24 checks whether anything was actually read. If not, we jump out of the buffered read loop to the second part of the program in line 25.

Now we are ready to scan the buffer. The symbol *\n* (line feed) is loaded into AL on line 26, and in line 27 this value is scanned for by REP SCASB and compared to the symbols in the buffer. There are two ways to exit the loop: either when CX hits zero or when a scanned symbol is a new line character. If the zero flag is set, then the last symbol scanned was a *\n* and the file position of the current symbol (one after the new line), is to be stored in the *linh* array. The count is then incremented and the file position is computed from BX and the number of characters still available is in CX (lines 29 through 31). Lines 32 through 34 perform the actual store, but since STOS assumes that the destination is in DI instead of in SI, these registers are exchanged before and after the STOS. Lines 35 through 37 check whether more data is available in the buffer, and jump according to the value of CX.

When the end of the file is reached, we have a complete list of file positions of the heads of the lines. Because we started the *linh* array with a 0 word, we know that the first line started at address 0, and that the next line starts at the position given by *linh* + 2 etc. The size of line *n* can be found from the starting address of line *n* + 1 minus the start address of line *n*.

The aim of the rest of the program is to read the number of a line, to read that line into the buffer, and to output it by means of a write call. All the necessary information can be found in the *linh* array, whose *n*th entry contains the position of the start of line *n* in the file. If the line number requested is either 0 or out of range, the program exits by jumping to label 7.

This part of the program starts with a call to the *getnum* subroutine on line 46. This routine reads a line from standard input and stores it in the *linein* buffer, (on lines 95 through 103). Next, we prepare for the SSCANF call. Considering the reverse order of the arguments, we first push the address of *curlin*, which can hold an integer value, then the address of the integer format string *numfmt*, and finally the address of the buffer *linein* containing the number in decimal notation. The system subroutine SSCANF puts the binary value in *curlin* if possible. On failure, it returns a 0 in AX. This return value is tested on line 48; on failure, the program generates an error message through label 8.

If the *getnum* subroutine returns a valid integer in *curlin*, then we first copy it in BX. Next we test the value against the range in lines 49 through 53, generating an EXIT if the line number is out of range.

Then we must find the end of the selected line in the file and the number of bytes to be read, so we multiply BX by 2 with a left shift SHL. The file position of the intended line is copied to AX on line 55. The file position of the next line is in CX and will be used to compute the number of bytes in the current line.

To do a random read from a file, an lseek call is needed to set the file offset to the byte to be read next. The lseek is performed with respect to the start of the file,

```

#include "../syscalnr.h" ! 1    PUSH _EXIT          ! 43    PUSH buf           ! 85
bufsiz = 512      ! 2    PUSH _EXIT          ! 44    PUSH (fildes)       ! 86
.SECT .TEXT        ! 3    SYS                ! 45    PUSH _READ         ! 87
inbufst:          ! 4    3: CALL getnum     ! 46    SYS                ! 88
    MOV BP,SP      ! 5    CMP AX,0          ! 47    ADD SP,8          ! 89
    MOV DI,linein   ! 6    JLE 8f            ! 48    MOV CX,AX          ! 90
    PUSH _GETCHAR   ! 7    MOV BX,(curlin)   ! 49    ADD BX,CX          ! 91
1: SYS             ! 8    CMP BX,0          ! 50    MOV DI,buf          ! 92
    CMPB AL,'\n'    ! 9    JLE 7f            ! 51    RET               ! 93
    JL 9f            ! 10   CMP BX,(count)   ! 52
    JE 1f            ! 11   JG 7f             ! 53    getnum:           ! 94
    STOSB           ! 12   SHL BX,1         ! 54    MOV DI,linein     ! 95
    JMP 1b           ! 13   MOV AX,linh-2(BX) ! 55    PUSH _GETCHAR     ! 96
1: PUSH 0          ! 14   MOV CX,linh(BX)  ! 56    1: SYS             ! 97
    PUSH linein     ! 15   PUSH 0            ! 57    CMPB AL,'\n'      ! 98
    PUSH _OPEN       ! 16   PUSH 0            ! 58    JL 9b              ! 99
    SYS              ! 17   PUSH AX            ! 59    JE 1f              ! 100
    CMP AX,0         ! 18   PUSH (fildes)    ! 60    STOSB             ! 101
    JL 9f            ! 19    PUSH _LSEEK       ! 61    JMP 1b             ! 102
    MOV (fildes),AX  ! 20   SYS                ! 62    1: MOVB (DI),'0'  ! 103
    MOV SI,linh+2    ! 21   SUB CX,AX        ! 63    PUSH curlin       ! 104
    MOV BX,0          ! 22   PUSH CX            ! 64    PUSH numfmt       ! 105
1: CALL fillbuf    ! 23   PUSH buf           ! 65    PUSH linein       ! 106
    CMP CX,0         ! 24   PUSH (fildes)    ! 66    PUSH _SSCANF      ! 107
    JLE 3f            ! 25   PUSH _READ        ! 67    SYS                ! 108
2: MOVB AL,'\n'     ! 26   SYS                ! 68    ADD SP,10         ! 109
    REPNE SCASB     ! 27   ADD SP,4          ! 69    RET               ! 110
    JNE 1b            ! 28   PUSH 1             ! 70
    INC (count)      ! 29   PUSH _WRITE       ! 71    .SECT .DATA      ! 111
    MOV AX,BX         ! 30   SYS                ! 72    errmess:          ! 112
    SUB AX,CX         ! 31   ADD SP,14         ! 73    .ASCIZ "Open %s failed\n" ! 113
    XCHG SI,DI        ! 32   JMP 3b            ! 74    numfmt: .ASCIZ "%d" ! 114
    STOS              ! 33   8: PUSH scanerr   ! 75    scanerr:          ! 115
    XCHG SI,DI        ! 34   PUSH _PRINTF     ! 76    .ASCIZ "Type a number.\n" ! 116
    CMP CX,0          ! 35   SYS                ! 77    .ALIGN 2           ! 117
    JNE 2b            ! 36   ADD SP,4          ! 78    .SECT .BSS        ! 118
    JMP 1b            ! 37   JMP 3b            ! 79    linein: .SPACE 80  ! 119
9: MOV SP,BP        ! 38   7: PUSH 0          ! 80    fildes: .SPACE 2   ! 120
    PUSH linein       ! 39   PUSH _EXIT        ! 81    linh: .SPACE 8192  ! 121
    PUSH errmess      ! 40   SYS                ! 82    curlin: .SPACE 4  ! 122
    PUSH _PRINTF      ! 41    fillbuf:          ! 83    buf: .SPACE bufsiz+2 ! 123
    SYS              ! 42    PUSH bufsiz      ! 84    count: .SPACE 2   ! 124

```

**Figure C-19.** A program with buffered read and random file access.

so first an argument of 0 is pushed to indicate this on line 57. The next argument is the file offset. By definition, this argument is a long (i.e., 32-bit) integer, so we first push a 0 word and then the value of AX on lines 58 and 59 to form a 32-bit integer. Then the file descriptor and code for LSEEK are pushed and the call is made on line 62. The return value of LSEEK is the current position in the file and can be found in the DX : AX register combination. If the number fits into a machine word (which it will for files shorter than 65536 bytes), then AX contains the address, so subtracting this register from CX (line 63), yields the number of bytes to be read in order to bring the line into the buffer.

The rest of the program is easy. The line is read from the file on lines 64 through 68 and then it is written to standard output via file descriptor 1 on lines 70 through 72. Note that the count and the buffer value are still on the stack after the partial stack cleanup on line 69. Finally, on line 73, we reset the stack pointer completely and we are ready for the next step, so we jump back to label 3, and restart with another call to *getnum*.

### Acknowledgements

The assembler used in this appendix is part of the “Amsterdam Compiler Kit.” The full kit is available online at [www.cs.vu.nl/ack](http://www.cs.vu.nl/ack). We thank the people who were involved in the original design: Johan Stevenson, Hans Schaminee, and Hans de Vries. We are especially indebted to Ceriel Jacobs, who maintains this software package, and who has helped adapt it several times to meet the teaching requirements of our classes, and also to Elth Ogston for reading the manuscript and testing the examples and exercises.

We also want to thank Robbert van Renesse and Jan-Mark Wams, who designed tracers for the PDP-11 and the Motorola 68000, respectively. Many of their ideas are used in the design of this tracer. Moreover, we wish to thank the large group of teaching assistants and system operators who have assisted us during many assembly language programming courses over a period of many years.

### PROBLEMS

1. After the instruction MOV AX, 702 is executed, what are the decimal values for the contents of AH and AL?
2. The CS register has the value 4. What is the range of absolute memory addresses for the code segment?
3. What is the highest memory address the 8088 can access?

4. Suppose that CS = 40, DS = 8000, and IP = 20.
  - a. What is the absolute address of the next instruction?
  - b. If MOV AX, (2) is executed, which memory word is loaded into AX?
5. A subroutine with three integer arguments is called following the calling sequence described in the text, that is, the caller pushes the arguments onto the stack in reverse order, then executes a CALL instruction. The callee then saves the old BP and sets the new BP to point to the saved old one. Then the stack pointer is decremented to allocate space for local variables. With these conventions, give the instruction needed to move the first argument into AX.
6. In Fig. C-1 the expression  $de - hw$  is used as an operand. This value is the difference of two labels. Might there be circumstances in which  $de + hw$  could be used as a valid operand? Discuss your answer.
7. Give the assembly code for computing the expression:

$$x = a + b + 2$$

8. A C function is called by

```
foobar(x, y);
```

Give the assembly code for making this call.

9. Write an assembly language program to accept input expressions consisting of an integer, an operator, and another integer and output the value of the expression. Allow the +, -, ×, and / operators..Pc

## **INDEX**

*This page intentionally left blank*

# INDEX

## A

Absolute path, 494  
Accelerated graphics port bus, 217  
Access control list, 501  
Access token, 501  
Accumulator, 19, 695  
Achieving high performance, 650, 650–652  
ACK (*see* Amsterdam Compiler Kit)  
Acknowledgment packet, 226  
ACL (*see* Access Control List)  
Acorn Archimedes computer, 45  
Active matrix display, 117  
Actual parameter, macro, 526  
Adder, 164–166  
    carry select, 165–166  
    full, 165–166  
    half, 164–165  
    ripller carry, 165  
Additive inverse, 390  
Address, 74, 694  
Address decoding, 233  
Address modes, 371  
Address space, 439  
Address-space identifier, 461

Addressing, 362, 371–386  
    8088, 701–704  
    ARM, 384  
    ATmega168, 384–385  
    based-indexed, 376  
    branch instructions, 380–381  
    Core i7, 382–384  
    direct, 372  
    displacement, 703  
    immediate, 372  
    indexed, 374  
    register, 372  
    register indirect, 373–374  
    stack, 376–379  
Addressing and memory, 8088, 699–704  
Addressing direct, 372  
Addressing modes, 384–385  
    ARM, 384  
    discussion, 385–386  
ADSL (*see* Asymmetric DSL)  
Advanced microcontroller bus architecture, 574  
Advanced programmable interrupt controller, 205  
Aggregate bandwidth, 647  
AGP bus (*see* Accelerated Graphics Port bus)  
Aiken, Howard, 16

- Algorithm, 8  
Allocation/rename unit, 327  
Alpha, 14, 26  
ALU (*see* Arithmetic Logic Unit)  
AMBA (*see* Advanced Microcontroller Bus Architecture)  
Amdahl's law, 649  
American Standard Code for Information Interchange, 138  
Amplitude modulation, 128  
Amsterdam compiler kit, 716  
Analytical engine, 15  
APIC (*see* Advanced Programmable Interrupt Controller)  
Apple II, 24  
Apple Lisa, 14  
Apple Macintosh, 24  
Apple Newton, 14, 45  
Application layer, 654  
Application programming interface, 487  
Application-level shared memory, 640–646  
Application-specific integrated circuit, 578  
Architecture, computer, 8, 60  
Arduino, 34  
Argument, 711  
Arithmetic circuit, 163–166  
Arithmetic logic unit, 6, 56–57, 166–167  
ARM,  
  addressing, 384  
  bi-endianess, 354  
  data types, 361  
  instruction formats, 368–370  
  instructions, 400–402  
  introduction, 45–47  
  ISA level, 354–356  
  microarchitecture of OMAP4430, 329–334  
  virtual memory, 460–462  
ARM (*see also* OMAP4430)  
ARM addressing modes, 385  
ARM data types, 361  
ARM instruction formats, 368–370  
ARM instructions, 400–402  
ARM v7, 346–347  
ARM virtual memory, 460–462  
As88, 716–720  
ASCII (*see* American Standard Code for Information Interchange)  
ASIC (*see* Application Specific Integrated Circuit)  
Assembler, 7, 518, 692, 715–721  
  8088, 715  
Assembler directive, 522  
Assembler pass one, 530–534  
Assembler pass two, 534–536  
Assembly language, 518–529, 691–745  
  example, 693–694  
Assembly-language level, 517–550  
Assembly-language macro, 524–529  
Assembly-language program, 716  
Assembly-language statement, 520–521  
Assembly process, 529–536  
Asserted signal, 178  
Associative memory, 454, 535  
Asymmetric DSL, 130  
Asynchronous bus, 194–196  
AT attachment disk, 92  
ATA packet interface, 92  
ATA-3 disk, 92  
Atanasoff, John, 16  
ATAPI-4 disk, 92  
ATAPI-5 disk, 92  
ATmega168, 212–214  
  ISA level, 356–358  
  addressing modes, 384–385  
  data types, 361–362  
  instruction formats, 370–371  
  instructions, 402  
  microarchitecture, 334–336  
Atmel ATmega168 (*see* ATmega168)  
Attraction memory, 615  
Auto decrement addressing, 704  
Auto increment addressing, 704  
Auxiliary carry flag, 708  
AVR, 47–49

**B**

- Babbage, Charles, 14–15  
Baby feeding algorithm, 445  
Backward compatibility, 344  
Ball grid array, 211  
Bank, 206  
Banking, Core i7, 328  
Bardeen, John, 19  
Base, 148  
Base pointer, 698  
Base register, 695  
Based-indexed addressing, 376

- Basic block, 320  
Basic input/output system, 91  
Batch system, 10–11  
Baud, 129  
Bayer filter, 135  
BCD (*see* Binary Coded Decimal)  
Benchmark, 293  
Best-fit algorithm, 454  
Bi-endian processor, 354  
Big endian memory, 76–78  
Binary, 671  
Binary arithmetic, 678  
Binary coded decimal, 74, 702  
Binary number,  
    addition of, 678  
    conversion between radices, 673–675  
    negative, 675–677  
Binary program, 692  
Binary search, 536  
Binding time, 542–545  
BIOS (*see* Basic Input Output System)  
Bipolar transistor, 150  
Bisection bandwidth, 618  
Bit, 74, 672  
Bit map, 360  
Bit slice, 166  
Block cache, 484  
Block started by symbol, 717  
Blocking network, 606  
Blu-ray disc, 108  
BlueGene, 622–626  
BlueGene/L, 622  
BlueGene/P, 622–626  
BlueGene/P vs. Red Storm, 629–631  
Boole, George,  
Boolean algebra, 147–157  
Boolean function, 152–153  
Branch history shift register, 314  
Branch instructions, addressing, 380–381  
Branch prediction, 310–315, 315  
    dynamic, 312–315  
Branch predictor, Core i7, 326  
Branch target buffer, 327  
Brattain, Walter, 19  
Breakpoint, 724  
Broadband, 129  
BSS section, 717  
BTB (*see* Branch Target Buffer)  
Bubblejet printer, 126  
Bucket, 536  
Buffered message passing, 637  
Bundle, Itanium 2, 425  
Burroughs B5000, 14, 21  
Bus, 20, 55, 108–112, 187–201, 214–232  
    ACP, 217  
    asynchronous, 194–196  
    Core i7, 206–209  
    EISA, 111  
    ISA, 110  
    multiplexed, 191  
    PCI, 111–112  
    PCIe, 111–112  
    synchronous, 191–194  
    system, 188  
Bus arbiter, 110  
Bus arbitration, 196–198  
Bus clocking, 191–194  
Bus cycle, 191  
Bus driver, 189  
Bus grant, 196  
Bus master, 189  
Bus operation, 198–201  
Bus protocol, 188  
Bus receiver, 189  
Bus skew, 112, 191  
Bus slave, 189  
Bus transceiver, 189  
Bus width, 190–191  
Busy waiting, 395  
Byron, Lord, 15  
Byte, 75, 347  
Byte instruction, 701  
Byte ordering, 76–78  
Byte register, 701

## C

- Cache, 304–310  
    direct-mapped, 306–308  
    level 2, 305  
    MESI protocol, 601–603  
    set-associative, 308–310  
    snooping, 599–606  
    split, 84, 305  
    unified, 84  
    update strategy, 600  
    write back, 310

- Cache (*continued*)  
     write-deferred, 310  
     write-through, 310, 599  
     write-allocate, 310  
 Cache coherence, 599  
 Cache coherence protocol, 599  
 Cache coherent NUMA, 607–608  
 Cache consistency, 599  
 Cache hit, 307  
 Cache invalidate strategy, 600  
 Cache line, 84, 306, 599  
 Cache memory, 41, 82–85  
 Cache miss, 307  
 Cache only memory access, 592, 614–616  
 Call gate, 460  
 Camera, digital, 135–137  
 Capacitor, 114  
 Carrier, 128  
 Carry select adder, 165–166  
 Catamount, 629  
 Cathode-ray tube, 115  
 CC-NUMA (*see* Cache Coherent NUMA)  
 CCD (*see* Charge-Coupled Device)  
 CD-R (*see* CD Recordable)  
 CD-Recordable, 103–105  
 CD-Rewritable, 105  
 CD-ROM (*see* Compact Disc-Read Only Memory)  
 CD-ROM track, 105  
 CD-ROM XA, 105  
 CDC (*see* Control Data Corporation)  
 CDC 6600, 14, 21  
 Celeron, 42  
 Cell, memory, 74  
 Central processing unit, 19, 55–73  
 Character code, 137–142  
 Charge-coupled device, 135  
 Checkerboarding, 453  
 Checksum generation, 581  
 Checksum verification, 580  
 Child process, 503  
 Chip, 158  
 Chipset, 205  
 Circuit,  
     arithmetic, 163–166  
     combinational, 159–163  
 Circuit equivalence, 153–158  
 CISC (*see* Complex Instruction Set Computer)  
 Clock, 168–169  
 Clock cycle time, 168  
 Clocked D latch, 171–172  
 Clocked SR latch, 170–171  
 Clone, 24  
 Closure, 670  
 Cloud computing, 38  
 Cluster, 37, 502, 593  
 Cluster computing, 631–636  
 Cluster of workstations, 593  
 CMYK printer, 125  
 Coarse-grained multithreading, 563  
 COBOL, 39, 359  
 Code generation, 716  
 Code page, 138  
 Code point, 139  
 Code segment, 694  
 Codesign, 27  
 Codeword, 78  
 Collective layer, 653  
 Collector, 148  
 Color gamut, 125  
 Color palette, 118  
 Color printing, 124–127  
 COLOSSUS, 14, 16–17  
 COMA (*see* Cache Only Memory Access computer)  
 COMA multiprocessor, 614  
 Combinational circuit, 159–163  
 Committed page, 490  
 Commodity off-the-shelf cluster, 37  
 Communicator, 637  
 Compact disc-read only memory, 101–103  
 Comparator, 162–163  
 Comparison instructions, 390–392  
 Comparison of instruction sets, 402  
 Compiler, 7, 518  
 Completeness property, 153  
 Complex instruction set computer, 62  
 Compute unified device architecture, 583  
 Computer,  
     data parallel, 70–72  
     disposable, 31–33  
     game, 35–36  
 Computer architecture, 8  
     milestones, 13–28  
 Computer center, 23  
 Computer zoo, 28–39  
 Condition code, 350  
 Condition code register, 699  
 Condition variable, 506  
 Conditional branch instructions, 390–392  
 Conditional execution, 427  
 Conditional jump, 710

- Consistency,  
cache, 599  
processor, 595–596  
release, 597–598  
sequential, 594–595  
strict, 594  
weak, 596–597, 597–598
- Consistency model, 594
- Constant pool, 260
- Consumer, 474
- Control Data Corporation, 21
- Control signal, 249
- Control store, 61, 253
- Controller, 109
- Conversion between radices, 673–675
- Coprocessor, 574–586
- Copy on write, 489
- Core, 25, 568
- Core 2 duo, 42
- Core dump, 10
- Core i7,  
addressing, 382–384  
addressing modes, 382–384  
banking, 328  
branch predictor, 326  
data types, 360  
hyperthreading, 564–568  
introduction, 35–47  
instruction formats, 367–368  
instructions, 397–400  
ISA level, 351–353  
memory model, 347–348  
microarchitecture, 323–329  
multiprocessor, 569  
photo of the die, 43  
pinout, 204–208  
pipelining, 206–208  
reorder buffer, 327  
retirement unit, 328  
scheduler, 327  
virtual memory, 455–460  
virtualization, 464
- CoreConnect, 572
- Coroutine, 410–413
- Cortex A9, 329–334
- COTS (*see* Commodity Off The Shelf)
- Counter register, 697
- COW (*see* Collection of Workstations)
- CP/M, 24
- CPU (*see* Central Processing Unit)
- CPU chip, 185–187
- CPU organization, 56
- Cray, Seymour, 21
- CRAY-1, 14
- CRC (*see* Cyclic Redundancy Check)
- CRC (*see* Cyclic Redundancy Code)
- Critical section, 509
- Crossbar switch, 604
- Crosspoint, 604
- CRT (*see* Cathode Ray Tube)
- Cryptography,  
public-key, 585  
symmetric-key, 585
- Cryptoprocessor, 585–586
- Cube network, 620
- CUDA (*see* Compute Unified Device Architecture)
- Cycle stealing, 110, 397
- Cyclic redundancy check, 226
- Cylinder, disk, 89
- D**
- D latch, 171–172
- Daisy chaining, 196
- Data cache, OMAP4430, 331
- Data center, 37
- Data movement instructions, 386–387
- Data parallel computer, 70–72
- Data path, 6, 56, 244–250  
Mic-1, 254  
Mic-2, 292  
Mic-3, 296  
Mic-4, 301
- Data path cycle, 58
- Data path timing, 247–249
- Data register, 697
- Data section, 717
- Data segment, 709
- Data type, nonnumeric, 359–360  
numeric, 358–359
- Data types, 358–362  
ARM, 360  
ATmega168, 361–362  
Core i7, 360
- DDR (*see* Double Data Rate RAM)
- De Morgan’s law 156

- DEC (*see* Digital Equipment Corporation)  
 DEC Alpha, 14, 26  
 DEC PDP-1, 14, 20  
 DEC PDP-11, 14  
 DEC PDP-8, 14, 20  
 DEC VAX, 14, 61  
 Decimal number, 718  
 Decoder, 161–162  
 Decoding unit, 300  
 Degree, node, 618  
 Delay slot, 311  
 Demand paging, 443–446, 445  
 Demultiplexer, 161  
 Denormalized number, 687  
 Dependence, 297, 318  
 Design principles, 63–65  
 Destination determination, 580  
 Destination index, 698  
 Destination operand, 701  
 Device level, 5  
 Device register bus, 573  
 Diameter, network 618  
 Dibit encoding, 129  
 Difference engine, 15  
 Digital camera, 135–137  
 Digital Equipment Corporation, 14, 19–20  
 Digital logic circuits, 158–169  
 Digital logic level, 5  
   buses, 185–201, 214–232  
   circuits, 158–169  
   CPU chips, 201–214  
   gates, 147–157  
 I/O interfacing, 232–235  
 memory, 169–185  
 PCI bus, 215–223  
 PCI express bus, 223–228  
 Digital subscriber line, 129–132, 130  
 Digital subscriber line access multiplexer, 132  
 Digital versatile disk, 106–108  
 Digital video disk, 106–108  
 Dimensionality, 619  
 DIMM (*see* Dual Inline Memory Module)  
 DIP (*see* Dual Inline Package)  
 Direct addressing, 372, 702  
 Direct memory access, 109, 396  
 Direct-mapped cache, 306  
 Direction flag, 704  
 Directory, 471  
 Directory management instructions, 471  
 Directory-based multiprocessor, 608  
 Disambiguation, 329  
 Disk,  
   ATAPI, 92  
   CD-ROM, 101–103  
   DVD, 106–108  
   IDE, 91–92  
   magnetic, 87–97  
   optical, 99–108  
   RAID, 94–97  
   SCSI, 92–94  
   SSD, 97–99  
   Winchester, 111  
 Disk controller, 90  
 Diskette, 89  
 Disposable computer, 31–33  
 Distributed memory system, 587  
 Distributed shared memory, 589–591, 640–642  
 DLL (*see* Dynamic Link Library)  
 DMA (*see* Direct Memory Access)  
 Dot, 716  
 Dots per inch, 126  
 Double data rate RAM, 181  
 Double indirect block, 497  
 Double integer, 697  
 Double torus, 620  
 Double-precision number, 358  
 DPI (*see* Dots Per Inch)  
 DRAM (*see* Dynamic RAM)  
 DSL (*see* Digital Subscriber Line)  
 DSLAM (*see* Digital Subscriber Line Access Multiplexer)  
 DSM (*see* Distributed Shared Memory)  
 Dual, 155  
 Dual inline memory module, 85  
 Dual inline package, 158  
 DVD (*see* Digital Video Disk)  
 Dyadic operations, 387–388  
 Dye sublimation printer, 127  
 Dye-based ink, 126  
 Dynamic branch prediction, 312–315  
 Dynamic link library, 547  
 Dynamic linking, 545–549  
   MULTICS, 545–547  
   UNIX, 549  
   Windows, 547–549  
 Dynamic RAM, 181  
 Dynamic voltage scaling, 209

**E**

Eckert, J. Presper, 17  
ECL (*see* Emitter-Coupled Logic)  
Edge-triggered flip-flop, 172  
EDO (*see* Extended Data Output)  
EDSAC, 14  
EDVAC, 17  
EEPROM (*see* Electrically Erasable PROM)  
Effective address, 704  
Egress processing, 580  
EHCI (*see* Enhanced Host Controller Interface)  
EIDE (*see* Extended IDE)  
EISA bus (*see* Extended ISA bus)  
Electrically erasable PROM, 182–184  
Electronic discrete variable automatic computer, 17  
Electronic numerical integrator and computer, 17  
Emitter, 148  
Emitter-coupled logic, 150  
Emulation, 22  
Encoding, 8b/10b, 226  
Endian memory,  
    big, 76–78  
    little, 76–78  
Enhanced host controller interface, 231  
ENIAC, 14, 17  
ENIGMA, 16  
Entry point, 542  
EPIC (*see* Explicitly Parallel Instruction Computing)  
EPROM (*see* Erasable PROM)  
EPT (*see* Extended Page Table)  
Erasable PROM, 182  
Error-correcting code, 78–82  
Escape code, 367  
Estridge, Philip, 24  
Ethernet, 575  
Event, 509  
Evolution of multilevel machines, 8–13  
Example computer families, 39  
Example programs, 8088, 726–745  
Excess notation, 676  
Executable binary file, 716  
Executable binary program, 518, 537  
Executive, NTOS, 487  
Expanding opcode, 365–367  
Explicit linking, 548  
Explicitly parallel instruction computing, 423  
Exponent, 682  
Extended data output, 181  
Extended IDE disk, 91

Extended ISA bus, 111  
Extended page table, 464  
External fragmentation, 453  
External reference, 539  
External symbol, 542  
Extra segment, 709

**F**

Fabric layer, 653  
False sharing, 641  
Fanout, 618  
Far call, 710  
Far jump, 8088, 709  
Fast page mode memory, 181  
FAT (*see* File Allocation Table)  
Fat tree, 620  
Fermi GPU, 70–71, 582–585  
Fetch-decode-execute cycle, 58, 244  
Field extraction, 580  
Field-programmable gate array, 25, 183–185, 578  
FIFO algorithm (*see* First-In First-Out algorithm)  
Fifth-generation computers, 26–28  
Fifth-generation project, Japanese, 26  
File, 465–467  
File allocation table, 498  
File descriptor, 492, 713  
File index, 468  
Filter, 494  
Fine-grained multithreading, 562  
Finite state machine, 290, 314  
Finite-precision numbers, 669–671  
Firewall, 576  
First-fit algorithm, 454  
First pass, assembler, 716  
First-generation computers, 16–19  
First-in, first-out algorithm, 447  
Flags register, 350, 699  
Flash memory, 183  
Flat panel display, 115–117  
Flip-flop, 172–174  
Floating-point number, 681–688  
Floppy disk, 89  
Flow control, 227  
Flow of control, 403–417, 404–417  
    branches, 405–406  
    routines, 410–413

- Flow of control (*continued*)  
 interrupts, 414–417  
 procedures, 406–410  
 sequential, 405–406  
 traps, 413
- Flynn’s taxonomy, 591–593
- Formal parameter, macro, 526
- Forrester, Jay, 19
- FORTRAN, 10, 392
- FORTRAN Monitor System, 10
- Forward reference problem, 529
- Fourth-generation computers, 23–26
- FPGA (*see* Field Programmable Gate Array)
- FPGA (*see* Field-Programmable Gate Array)
- FPM (*see* Fast Page Mode)
- Fraction, 682
- Fragmentation, 448–449  
 external, 453  
 internal, 448
- Fragmentation and reassembly, 581
- Frame, 101  
 local variable, 258–260
- Frame pointer, 353
- Free list, 469
- Free page, 490
- Frequency modulation, 128
- Frequency shift keying, 128
- FSM (*see* Finite State Machine)
- Full adder, 165–166
- Full handshake, 195
- Full interconnect, 620
- Full resource sharing, 566
- Full-duplex line, 129
- Function, 392
- G**
- Game computer, 35–36
- Game controller, 120–122
- Gamut, color, 125
- Gate, 5, 148–150
- Gate delay, 159
- GDT (*see* Global Descriptor Table)
- General registers, 695–697
- General-purpose CPU, 584
- Ghosting, 114
- Global descriptor table, 455
- Global label, 717
- Goldstine, Herman, 18
- Google cluster, 632–636
- GPGPU (*see* General Purpose GPU)
- GPU (*see* Graphics Processing Unit)
- Graphical user interface, 24, 485
- Graphics processing unit, 70–72, 582–585
- Graphics processor, 582–585
- Green book, 103
- Grid, 620
- Grid computing, 652–655
- GridPad, 14
- Guest operating system, 464
- GUI (*see* Graphical User Interface)
- H**
- H register 245–246
- Half adder, 164–165
- Half-duplex line, 129
- Halftone screen frequency, 124
- Halftoning, 124
- Hamming, Richard, 30
- Hamming distance, 78
- Handle, 488
- Hard disk, 87–97
- Hardware, 8  
 equivalence with software, 8
- Hardware abstraction layer, 486
- Hardware DSM, 607
- Hardware virtualization, 463–464
- Harvard architecture, 84
- Harvard Mark I, 16
- Hash coding, 536
- Hashing, 536
- Hawkins, Jeff, 27
- Hazard, 297
- Headend, 132
- Header, PCI packet, 225
- Header management, 581
- Headless workstation, 631
- Hello world example, 8088, 726
- Heterogeneous multiprocessor, 570–574
- Hexadecimal, 671
- Hexadecimal number, 718
- High Sierra, 103
- High-level language, 7

- History,  
  1642–1945, 13–16  
  1945–1955, 16–19  
  1955–1965, 19–21  
  1965–1980, 21–23  
  ARM, 45–47  
  AVR, 47–49  
  Intel, 39–45  
Hit ratio, 84  
Hoagland, Al, 30  
Hoff, Ted, 40  
Hoisting, 321  
Homogeneous multiprocessors, 568–569  
Host library, 549  
HTTP (*see* HyperText Transfer Protocol)  
Hypercube, 621  
Hypertext transfer protocol, 576  
Hyperthreading, Core i7, 564–568  
Hypervisor, 463
- I**
- I-node, 496  
I/O (*see* Input/Output)  
I/O interface, 232–236  
IA-32 architecture, 351  
IA-32 flaws, 421–423  
IA-64, 420–429, 423–429  
  bundle, 425  
  EPIC model, 423  
  instruction scheduling, 424–426  
  speculative loads, 429  
IAS, 14  
IAS machine, 18  
IBM 1401, 14, 20, 21–22  
IBM 360, 14, 22–23  
IBM 701, 19  
IBM 704, 19  
IBM 709, 10  
IBM 7094, 14, 20–21, 21–22  
IBM 801, 62  
IBM CoreConnect, 572  
IBM Corporation, 19–21  
IBM PC, 14, 24–25  
IBM POWER4, 14, 26  
IBM PS/2, 41  
IBM RS6000, 14
- IBM Simon, 14  
IC (*see* Integrated Circuit)  
IDE disk, 91–92  
IEEE Floating-point standard 754, 684–688  
IFU (*see* Instruction Fetch Unit)  
IJVM, 258–282  
  compiling Java to, 265–267  
IJVM (*see* Integer JVM)  
IJVM implementation, 271–282  
IJVM instruction set, 262–265  
IJVM memory model, 260  
ILC (*see* Instruction Location Counter)  
ILLIAC, 17, 70  
Immediate addressing, 372, 704  
Immediate file, 503  
Immediate operand, 372  
Implicit linking, 548  
Implied addressing, 704  
Import library, 548  
Imprecise interrupt, 318  
Improving network performance, 582  
Improving performance, 650–652  
Index register, 698  
Indexed addressing, 374–376  
Indexed color, 118  
Indirect block, 497  
Industry standard architecture bus, 110  
Infix, 376  
Informative section, 346  
Ingress processing, 580  
Initiator, PCI bus, 218  
Ink,  
  dye-based, 126  
  pigment-based, 126  
  solid, 126  
Inkjet printer, 125–126  
Input/output, 108–142  
Input/output device,  
  color printer, 124–125  
  digital camera, 135–137  
  flat panel display, 115–118  
  game controller, 120–122  
  inkjet printer, 125–126  
  keyboard, 113  
  laser printer, 122–124  
  mice, 118–120  
  modem, 127–129  
  specialty printers, 126–127  
  telecommunication equipment, 127–132  
  touch screen, 113–115

- Input/output instructions, 394–397  
 Instruction,  
   ARM 400–402  
   ATmega168, 402  
   comparison, 390–392  
   conditional branch, 390–392  
   Core i7, 397–400  
   data movement, 386–387  
   dyadic, 387–388  
   input/output, 394–397  
   ISA level, 351  
   loop control, 393–394  
   monadic, 388–390  
   predicated, 370  
   procedure call, 392–393  
 Instruction execution, 58–62  
 Instruction fetch unit, 288–291  
 Instruction format, 362–371  
   ARM, 368–370  
   ATmega168, 370–371  
   Core i7, 367–368  
   design criteria, 362–365  
 Instruction group, Itanium 2, 424  
 Instruction issue unit, OMAP4430, 330  
 Instruction location counter, 530  
 Instruction pointer, 694, 698  
 Instruction register, 56  
 Instruction scheduling, Itanium 2, 424  
 Instruction set, 8088, 705–715  
 Instruction set architecture, 691  
 Instruction set architecture level, 343–430  
 Instruction types, 386–403  
 Instruction-level parallelism, 65–69, 555–561  
 Instructions,  
   ARM, 400–402  
   ATmega168, 403  
   Core i7, 397–400  
 Integer JVM, 258–267  
   compiling Java to, 265–267  
   implementation, 267–282  
   instruction set, 262–265  
   memory model, 260–261  
   stack, 258–260  
 Integrated circuit, 158–159  
 Integrated drive electronics, 91  
 Intel 386, 14  
 Intel 4004, 40  
 Intel 80286, 41  
 Intel 80386, 41  
 Intel 80486, 41  
   Intel 8080, 14  
   Intel 8086, 41  
   Intel 8088, 694–704  
     addressing, 699–704  
     example programs, 726–745  
     far jump, 709  
     near jump, 709  
     segments, 700  
   Intel 8088 instruction set, 705–715  
   Intel Core i7 (*see* Core i7)  
   Intel Xeon, 43  
 Interconnection network, 617–621  
   bisection bandwidth, 618  
   topology, 618  
 Interfacing, 232–236  
 Interleaved memory, 606  
 Internal fragmentation, 448  
 Internet over cable, 132–135  
 Internet protocol, 576  
 Internet service provider, 576  
 Interpretation, 2  
 Interpreter, 2, 58–59, 693  
 Interrupt, 109, 414–417  
   imprecise, 318  
   precise, 318  
   transparent, 415  
 Interrupt handler, 109  
 Interrupt vector, 200, 414  
 Intersector gap, 87  
 Introduction, 715  
 Invalidate strategy, 600  
 Inversion bubble, 149  
 Inverter, 149  
 Inverting buffer, 177  
 Invisible computer, 26–28  
 IP header, 576  
 IP protocol, 576  
 IR (*see* Instruction Register)  
 Iron Oxide Valley, 20  
 ISA bus (*see* Industry Standard Architecture bus)  
 ISA level, 6  
   addressing, 371–386  
   ARM, 354–356  
   ATmega168, 356–358  
   Core i7, 351–353  
   data types, 358–362  
   flow of control, 404–417  
   instruction formats, 362–371  
   instruction types, 386–404  
   overview, 345–358

ISA level properties, 345–347  
ISP (*see* Internet Service Provider)  
Itanium 2, 420–429

## J

Java virtual machine, 260  
Jobs, Steve, 24  
Johnniac, 17  
Joint Photographic Experts Group, 137  
Joint test action group, 206  
JPEG (*see* Joint Photographic Experts Group)  
JTAG (*see* Joint Test Action Group)

## K

Kernel mode, 347  
Key, record, 468  
Keyboard, 113  
Kilby, Jack, 21  
Kildall, Gary, 24  
Kinect controller, 122

## L

Label, 692  
    assembly language, 717  
LAN (*see* Local Area Network)  
Land, 100  
Land grid array, 158  
Lane, PCI Express, 112, 226  
Language, 1  
Large memory model, 720  
Laser printer, 122–134  
Latch, 169–172  
    D, 171–172  
Latency, 67  
Latin-1, 138  
Layer, 3  
LBA (*see* Logical Block Addressing)  
LCD (*see* Liquid Crystal Display)

LDT (*see* Local Descriptor Table)  
Least recently used algorithm, 309, 446  
LED (*see* Light Emitting Diode)  
Left value, 701  
Leibniz, Gottfried Wilhelm von, 13  
Level, 3  
    assembly-language, 517–550  
    device, 5  
    digital logic, 147–237  
    ISA, 6, 343–430  
    microarchitecture, 6, 243–338  
    operating-system machine, 7, 437–510  
Level 2 cache, 305  
Level-triggered latch, 172  
LGA (*see* Land Grid Array)  
Light emitting diode, 119  
Linda, 642–644  
Linear address, 457  
Lines per inch, 124  
Link, file, 494  
Link layer, PCI Express, 226  
Linkage editor, 536  
Linkage segment, 545  
Linker, 536, 716  
    tasks performed, 538–541  
Linking, 536–549  
    binding time, 542–545  
    dynamic, 545–549  
    MULTICS, 545–547  
    UNIX, 549  
    Windows, 547–548  
Linking loader, 536  
Liquid crystal display, 115  
Lisa, 14  
Literal, 532  
Little endian, 702  
Little endian memory, 76–78  
Load/store architecture, 356  
Loading, 536–549  
Local descriptor table, 455  
Local label, 717  
Local loop, 130  
Local variable frame, 258–260  
Local-area network, 575  
Locality principle, 83, 446  
Location, 74  
Location counter, 716  
Logical block addressing, 91  
Logical record, 466  
Long integer, 697

- Lookup table, 183–185  
 Loop control instructions, 393–394  
 Loosely coupled system, 73, 554  
 Lovelace, Ada, 15  
 LPI (*see* Lines Per Inch)  
 LRU (*see* Least Recently Used algorithm)  
 LRU algorithm, (*see* Least Recently Used algorithm)  
 LRU algorithm (*see* Least Recently Used algorithm)  
 LUT (*see* LookUp Table)
- M**
- M.I.T.  
   TX-0, 19  
   TX-2, 19  
 Machine language, 1, 691  
 Macintosh, Apple, 24  
 Macro, assembly language, 524–529  
 Macro call, 525  
 Macro definition, 524  
 Macro expansion, 525  
 Macro parameter, 526–527  
 Macro processor, implementation, 528–529  
 Macroarchitecture, 258  
 Magnetic disk, 87–97  
 Mainframe, 38–39  
 MAL (*see* Micro Assembly Language)  
 MANIAC, 17  
 Mantissa, 682  
 MAR (*see* Memory Address Register)  
 Mark I, 14  
 Mask, 387  
 MASM (*see* Microsoft ASseMbler)  
 Massively parallel processor, 593, 621–631
  - BlueGene, 622–626
  - Red Storm, 626–631
 Master, bus, 189  
 Master file table, 502  
 Masuoka, Fujio, 98  
 Mauchley, John, 16  
 MDR (*see* Memory Data Register)  
 Memory, 73–85, 169–185
  - associative, 454, 535
  - attraction, 615
  - big endian, 76–78
  - cache, 41, 82–85
  - cache only access, 614–616
  - compact disc-read only, 101–103
  - fast page mode, 181
  - flash, 183
  - interleaved, 606
  - little endian, 76–78
  - nonvolatile, 182–183
  - random access, 73–85, 180–183
  - read only, 182
  - secondary, 86–108
  - virtual, 438–462
 Memory address, 74–76  
 Memory address register, 249–250  
 Memory and addressing, 8088, 699–704  
 Memory chip, 178–180  
 Memory data register, 249–250  
 Memory hierarchy, 86–87  
 Memory management unit, 442  
 Memory map, 440  
 Memory model, 347–349, 720
  - Core i7, 347–348
 Memory operation, 249–250  
 Memory organization, 174–178
  - 8088, 699–701
 Memory packaging, 85  
 Memory semantics, 593–598
  - cache, 599
  - processor, 595–596
  - release, 597–598
  - sequential, 594–595
  - strict, 594
  - weak, 596–597
 Memory-mapped I/O, 234  
 Mesh, 620  
 MESI protocol, 601–603  
 Message queue, 504  
 Message-passing interface, 637–639  
 Message-passing multicomputer, 616–652  
 Metal oxide semiconductor, 150  
 Method, 392  
 Method area, 260  
 Metric units, 49–50  
 MFT (*see* Master File Table)  
 Mic-1, 253, 253–258, 267–282  
 Mic-2, 291–293  
 Mic-3, 293–300  
 Mic-4, 300–303  
 Mickey, 120  
 Micro assembly language, 268–271  
 Micro-op cache, 325  
 Micro-operation, 300

- Microarchitecture,  
ARM, 329–324  
ATmega168, 334–336  
Core i7, 323–329  
Mic-1, 267–282  
Mic-2, 291–293  
Mic-3, 293–300  
Mic-4, 300–303  
OMAP4430, 329–334  
three-bus, 287
- Microarchitecture level, 6, 243–338  
design, 283–303  
examples, 323–336
- Microarchitecture of the Core i7 CPU, 323
- Microcode, 11–13
- Microcontroller, 33–35
- Microdrive, 137
- Microinstruction, 62, 251–258
- Microinstruction register, 255
- Microprogram, 6, 695
- Microprogram counter, 255
- Microprogramming, 9, 12–13
- Microsoft assembler, 520
- Microstep, 297
- Milestones in computer architecture, 13–28
- MIMD (*see* Multiple Instruction stream  
Multiple Data stream computer)
- Minislot, 134
- MIPS, 14
- MIPS computer, 62
- MIR (*see* MicroInstruction Register)
- Miss ratio, 84
- MMU (*see* Memory Management Unit)
- MMX (*see* MultiMedia eXtensions)
- Mnemonic, 692, 715
- Modem, 127–129
- Modulation, 128
- Modulator demodulator, 128
- Monadic operations, 388–390
- Monitor, 115–117
- Moore, Gordon, 28
- Moore’s law, 28
- MOS (*see* Metal Oxide Semiconductor)
- Motherboard, 108
- Motif, 485
- Motion Picture Experts Group, 571
- Motorola 68000, 61
- Mouse, 118–120
- MPC (*see* MicroProgram Counter)
- MPEG-2, 571
- MPI (*see* Message Passing Interface)
- MPP (*see* Massively Parallel Processor)
- MS-DOS, 25
- Multicomputer, 73, 587–591, 616–652, 621–631  
BlueGene, 622–626  
Google cluster, 632–636  
MPP, 593  
Red Storm, 626–631
- Multicomputer performance, 650–652
- Multicomputer software, 636–639, 636–646
- MULTICS (*see* MULTplexed Information and  
Computing Service)
- Multilevel machine, 5–13
- Multilevel machines, evolution, 8–13
- Multimedia extensions, 41
- Multiple instruction stream multiple data  
stream computer, 591–592
- Multiplexed bus, 191
- Multiplexed information and computing  
service, 454–455, 545–547
- Multiplexer, 159–161
- Multiprocessor, 72–73, 586–616  
COMA, 614–616  
Core i7, 569  
heterogeneous, 570–574  
NUMA, 606–610  
symmetric, 587, 598–606  
vs. multicomputer, 586–593
- Multiprogramming, 22
- Multisession CD-ROM, 105
- Multistage switching network, 604–606
- Multitouch screen, 114
- Mutex, 505
- Mutual capacitance, 114
- Myhrvold, Nathan, 29

## N

- N-way set associative cache, 308–310
- NaN (*see* Not a Number)
- Nathan’s law, 29
- NC-NUMA (*see* No Cache NUMA)
- Near call, 710
- Near jump, 8088, 709
- Negated signal, 178
- Negative binary numbers, 675–677
- Negative logic, 157

- NEON, 331  
 Network,  
   Ethernet, 575  
   local-area, 575  
   ring, 569  
   store-and-forward, 575  
   wide-area, 575  
 Network interface device, 131  
 Network of workstations, 593  
 Network processor, 574–582, 577–579, 578  
 Networking, introduction, 575–577  
 Newton, 14, 45  
 Nibble, 399  
 NID (*see* Network Interface Device)  
 No cache NUMA, 606  
 No remote memory access computer, 593  
 Nonblocking message passing, 637  
 Nonblocking network, 604–606  
 Noninverting buffer, 177  
 Nonnumeric data type, 359–360  
 Nonuniform memory access, 606–610  
 Nonuniform memory access computer, 592  
 Nonvolatile memory, 182–183  
 NORMA (*see* NO Remote Memory Access computer)  
 Normalized floating-point number, 684  
 Normative section, in standard, 346  
 Not a number, 688  
 Notation, Mic-1, 267–271  
 NOW (*see* Network Of Workstations)  
 Noyce, Robert, 21  
 NT File System, 498  
 NTFS (*see* NT File System)  
 NTOS executive, 487  
 NUMA (*see* NonUniform Memory Access computer)  
 NUMA multiprocessor, 606–610  
 Numeric data type, 358–359  
 Nvidia Fermi GPU, 582–585  
 Nvidia Tegra, 46
- OCP-IP (*see* Open Core Protocol-International Partnership)  
 Octal, 671  
 Octal number, 718  
 Octet, 75, 347  
 OGSA (*see* Open Grid Services Architecture)  
 OHCI (*see* Open Host Controller Interface)  
 OLED (*see* Organic Light Emitting Diode)  
 Olsen, Kenneth, 19  
 OMAP4430,  
   addressing, 384  
   bi-endianness, 354  
   data cache, 331  
   data types, 361  
   instruction formats, 368–370  
   instruction issue unit, 330  
   instructions, 400–402  
   internal organization, 209  
   introduction, 45–47  
   ISA level, 354–356  
   microarchitecture, 329–334  
   pinout, 211  
   pipeline, 331–334  
   store buffer, 331  
   virtual memory, 460–462  
 Omega network, 604–606  
 Omnibus, PDP-8, 20  
 On-chip multithreading, 562–568  
 On-chip parallelism, 554–574  
 One's complement number, 675  
 Opcode, 244  
 Open collector, 189  
 Open core protocol-international partnership, 574  
 Open grid services architecture, 654  
 Open host controller interface, 231  
 Operand stack, 259–260  
 Operating system, 437  
   CP/M, 24  
   history, 9–11  
   OS/2, 25  
   timesharing, 11  
   UNIX, 482–485, 488–489, 492–498, 503–506  
   Windows, 490–492  
   Windows 7, 485–488, 498–503, 506–509  
 Operating-system machine level, 7, 437–510  
 Operating-system macro, 11  
 Operation code, 244  
 Optical discs, 99–108  
 Orange book, 105  
 Orca, 644–646
- O**
- Object file, 716  
 Object module, 541–542  
 Object program, 518

- Organic light emitting diode, 117  
OS/2, 25  
Osborne-1, 14  
Out-of-order execution, 315–320  
Overlay, 439
- P**
- Packet, 575, 618  
  PCI, 225  
Packet classification, 580  
Packet processing, 580–581  
Packet processing engine, 579  
Packet switching, 575  
Page, 440  
Page directory, 457  
Page fault, 443  
Page frame, 442  
Page scanner, 607  
Page table, 440  
Page-replacement policy, 446–448  
Paging, 439–452  
  implementation, 441–443  
Paging algorithm, 446–448  
  FIFO, 447  
  LRU, 446  
Palm PDA, 27  
Parallel computer,  
  coprocessor, 574–586  
  grid, 652–655  
  multicomputer, 616–652  
  multiprocessor, 586–61  
  on-chip parallelism, 554–574  
Parallel computers,  
  performance, 646–652  
  taxonomy, 591–593  
Parallel input/output, 232  
Parallel processing, 471–480  
Parallel virtual machine, 637  
Parallelism, instruction-level, 65–69, 555–561  
  on-chip, 554–574  
Parameter, macro, 526–537  
Parent process, 503  
Parity bit, 79  
Parity flag, 708  
Partial address decoding, 235  
Partitioned resource sharing, 566
- Pascal, Blaise, 13  
Pass one, assembler, 530–534  
Pass two, assembler, 534–536  
Passive matrix display, 117  
Path, 494  
Path length, 283–291  
  reducing, 285–291  
Path selection, 580  
Payload, PCI packet, 225  
PC (*see* Program counter)  
PCI Bus (*see* Peripheral Component Interconnect Bus)  
PCI Express, 223–228  
  architecture, 224–225  
PCI Express bus, 111–112  
PCI Express protocol stack, 225–228  
PCIe (*see* PCI Express)  
PCIe bus (*see* PCI Express bus)  
PDA (*see* Personal Digital Assistant)  
PDP-1, 14, 20  
PDP-11, 14  
PDP-8, 14, 20  
Pentium, 25, 41  
Pentium 4, 44  
Perfect shuffle, 605  
Performance,  
  hardware metrics, 647–648  
  improving, 650–652  
  improving network, 582  
  software metrics, 648–650  
Performance metrics, 647–650  
  hardware, 647–648  
  software, 648–650  
Performance of parallel computers, 646–652  
Peripheral bus, 573  
Peripheral component interconnect bus,  
  111–112, 215–218  
  arbitration, 219–220  
  operation, 218–219  
  signals, 220–222  
  transactions, 222–223  
Perpendicular recording, 88  
Personal computer, 23–25, 36  
Personal digital assistant, 27  
Pervasive computing, 28  
PGA (*see* Pin Grid Array)  
Phase modulation, 128  
Physical address space, 440  
Physical layer, PCI Express, 226  
Pigment-based ink, 126  
Pin grid array, 158

- Pinout, 185  
 PIO (*see* Parallel Input/Output)  
 Pipe, 504  
 Pipeline, 65–67, 293  
   Mic-3, 298–300  
   Mic-4, 300–303  
   OMAP, 331–334  
   Pentium, 68  
   seven-stage, 300–303  
 Pipeline stall, 298  
 Pipelining, Core i7 memory, 206–208  
 Pit, 100  
 Pixel, 117  
 Plain old telephone service, 130  
 PlayStation 3, 35  
 Pointer, 373  
 Pointer register, 698–699  
 Poison bit, 322  
 Polish notation, 376–379  
 Portable Operating System-IX, 483  
 Position independent code, 545  
 Positive logic, 157  
 POSIX (*see* Portable Operating System-IX)  
 Postfix, 376  
 POTS (*see* Plain Old Telephone Service)  
 Power gating, 209  
 PPE (*see* Packet Processing Engine)  
 Preamble, 87  
 Precise interrupt, 318  
 Predicated instruction, 370  
 Predication, 426–428  
 Prefetch buffer, 65  
 Prefetching, 651  
 Prefix, 400  
 Prefix byte, 278, 367  
 Present/absent bit, 442  
 Print engine, 123  
 Printer, 122–127, 124–127  
   bubblejet, 126  
   CMYK, 125  
   dye-sublimination, 127  
   inkjet, 125–126  
   laser, 122–124  
   solid-ink, 126  
   specialty, 126–127  
   wax, 127  
 Procedure, 392, 406–410  
 Procedure call instructions, 392–393  
 Procedure epilog, 409  
 Procedure prolog, 409  
 Process creation, 473  
 Process management, 503, 506–509  
   UNIX, 503–506  
   Windows 7, 506–509  
 Process synchronization, 478–480  
 Processor, 55–73  
 Processor bandwidth, 67  
 Processor bus, 572  
 Processor consistency, 595–596  
 Processor cycle, 695  
 Processor-level parallelism, 69–73  
 Producer, 474  
 Program, 1  
 Program counter, 56, 694  
 Program status word, 350, 459  
 Programmable interconnect, 183  
 Programmable processing engine, 579  
 Programmable ROM, 182  
 Programmed I/O, 394  
 PROM (*see* Programmable ROM)  
 Protocol, 576  
   Bus, 188  
   IP, 576  
   MESI, 601–603  
   PCI Express, 225–228  
 Protocol processing engine, 579  
 Protocol stack, PCI Express, 225–228  
 Pseudoinstruction, 522–524, 692, 718  
 PSW (*see* Program Status Word)  
 Pthreads, 505  
 Public-key cryptography, 585  
 PVM (*see* Parallel Virtual Machine)

**Q**

- Queue management, 581  
 Queueing unit, 300

**R**

- Race condition, 473–478  
 Radio frequency identification, 31–33  
 Radix, 671  
 Radix number systems, 671–673

- RAID (*see* Redundant Array of Inexpensive Disks)  
RAM (*see* Random Access Memory)  
Random access memory, 73–85, 180–183  
    DDR, 181  
    dynamic RAM, 181  
    SDRAM, 181  
Ranging, 134  
RAW dependence, 297  
Read only memory, 182  
Read/write pointer, 714  
Real mode, 352  
Recursion, 392  
Recursive procedure, 406  
Red book, 100  
Red Storm, 626–631  
Red Storm vs. BlueGene/P, 629–631  
Reduced instruction set computer, 62  
    design principles, 63–65  
    vs. CISC  
Redundant array of inexpensive disks, 94–97  
Reed-Solomon code, 87  
Refresh, memory, 181  
Register, 5, 174, 694  
    PSW, 350  
Register addressing, 372  
Register displacement, 703  
Register indirect addressing, 373–374, 702  
Register renaming, 315–320  
Register with index, 703  
Register with index and displacement, 704  
Registers, 349–351  
    flags, 350  
Relative error, 683  
Relative path, 494  
Release consistency, 597–598  
Relocation constant, 541  
Relocation problem, 539  
Reorder buffer, Core i7, 327  
Replicated worker model, 644  
Reserved page, 490  
Resource layer, 653  
Retirement unit, Core i7, 328  
Reverse Polish notation, 376–379  
RFID (*see* Radio Frequency IDentification)  
Right-justified alignment, 387  
Ring network, 569, 620  
RISC (*see* Reduced Instruction Set Computer)  
RISC design principles, 63–65  
RISC vs. CISC, 62–63  
ROB (*see* ReOrder Buffer)  
ROM (*see* Read Only Memory)  
Root directory, 494  
Root hub, 229  
Rotational latency, 89  
Rounding, 683  
Route lookup, 581  
Router, 575
- S**
- Sandy Bridge, 323–329  
Saturated arithmetic, 559  
Scalable, 650  
Scalable design, 589  
Scale, index, base, byte 368, 383  
Scheduler, Core i7, 327  
Scheduling, multicomputer, 639–640  
Scoreboard, 316  
SCSI (*see* Small Computer System Interface)  
SCSI disk, 92  
SDRAM (*see* Synchronous DRAM)  
Seastar, 627  
Second pass, assembler, 716  
Second-generation computers, 19–21  
Secondary memory, 86–108  
Sector, 87  
Security descriptor, 501  
Security ID, 501  
Seek, disk, 89  
Segment, 450, 699  
    8088, 700  
Segment override, 715  
Segment register group, 699  
Segmentation, 449–455  
    implementation, 452–455  
Self-modifying, 374  
Semantics, memory, 593–598  
Semaphore, 478–480  
Sequencer, 253  
Sequential consistency, 594–595  
Sequential flow of control, 405–406  
Serial ATA, 92  
Server, 36  
Session, CD-ROM, 105  
Set-associative cache, 308–310  
Shard, 632  
Shared library, 549

- Shared-memory, application-level, 640–646  
 Shell, 485  
 Shifter, 163–164  
 Shockley, William, 19  
*SIB* (*see* Scale, Index, Base byte)  
*SID* (*see* Security ID)  
 Sign extension, 250  
 Signed magnitude number, 675  
 Significand, 686  
*SIMD* (*see* Single Instruction stream Multiple Data stream)  
*SIMM* (*see* Single Inline Memory Modules)  
 Simon smartphone, 27  
 Simple COMA, 615  
 Simplex line, 129  
 Simultaneous multithreading, 564  
 Single inline memory module, 85  
 Single instruction stream multiple data stream, 70, 583  
 Single large expensive disk, 94  
 Single-chip multiprocessors, 568–574  
 Skew, bus, 112  
 Slave, 189  
     bus, 189  
*SLED* (*see* Single Large Expensive Disk)  
*SM* (*see* Streaming Multiprocessor)  
 Small computer system interface, 92–94  
 Small memory model, 720  
 Small outline DIMM, 85  
 Smartphone, 27  
*SMP* (*see* Symmetric MultiProcessor)  
 Snoop, 202  
 Snooping cache, 599, 599–606  
 Snoopy cache, 599  
*SO-DIMM* (*see* Small Outline DIMM)  
 SoC, 208  
 Socket, 483  
 Software, 8  
 Software layer, 227  
 Software metric, 648  
 Solid-ink printer, 126  
 Solid-state disk, 97–99  
 Source index, 698  
 Source language, 517  
 Source operand, 701  
 SPARC, 14  
 Spatial locality, 305  
 Specialty printer, 126–127  
 Speculative execution, 320–323, 321  
 Speculative load, 429  
 Speed vs cost, 283–285  
 Split cache, 84, 305  
 Splitter, 131  
 SR latch, 169–170  
*SRAM* (*see* Static RAM)  
*SSD* (*see* Solid-State Disk)  
*SSE* (*see* Streaming SIMD Extensions)  
 Stack, 258–260  
     operand, 259  
 Stack addressing, 376–379  
 Stack frame, 698  
 Stack pointer, 698  
 Stage, pipeline, 65  
 Stale data, 599  
 Stall, 311  
 Stalling, 298  
 Standard error, 494  
 Standard input, 494  
 Standard output, 494  
 Star, 620  
 State, 244  
     finite state machine, 290  
 Statement, assembly language, 520–521  
 Static branch prediction, 315  
 Static RAM, 181  
 Stibitz, George, 16  
 Storage, 73  
 Store, 73  
 Store buffer, OMAP4430, 331  
 Store-and-forward network, 575  
 Store-to-load forwarding, 329  
 Stream, 484  
 Streaming multiprocessor, 582  
 Streaming SIMD extensions, 42  
 Strict consistency, 594  
 Striping, 95  
 Strobe, 171  
 Structured computer organization, 2–13  
 Subroutine, 392, 710  
 Sun Fire E25K, 610–614  
 Sun SPARC, 14  
 Supercomputer, 21  
 Supercomputers, 39  
 Superscalar architecture, 67–69, 68–69  
 Superuser, 496  
 Supervisor call, 11  
 Switching algebra, 150  
 Switching network, multistage, 604–606  
 Symbol table, 530, 535–536, 716  
 Symbolic name, 692  
 Symmetric key cryptography, 585

Symmetric multiprocessor, 587, 598–606  
Synchronous bus, 191–194  
Synchronous DRAM, 181  
Synchronous memory interface, 208  
Synchronous message passing, 636  
System bus, 188  
System call, 11, 437, 712–715  
System-on-a-chip, 208  
Systems programmer, 7

## T

Target, 218  
Target language, 517  
Target library, 549  
Task bag, 644  
TAT-12/13 cable, 30  
Taxonomy of parallel computers, 591–593  
TCP (*see* Transmission Control Protocol)  
TCP header, 576  
Tegra, Nvidia, 46  
Telco, 129  
Telecommunications equipment, 127–135  
Telephone company, 129  
Template, 643  
Temporal locality, 306  
Terminal, 113–118  
Texas Instruments OMAP4430 (*see* OMAP4430)  
Text section, 717  
TFT display, 117  
Thermal printer, 127  
Thermal throttling, 206  
Thin film transistor, 117  
Third-generation computers, 21–23  
Thrashing, 448  
Thread, 505  
Three-bus architecture, 287  
Threshold sharing, 567  
TI OMAP4430 (*see* OMAP4430)  
Tightly coupled system, 73, 554  
Timesharing system, 11  
Tiny memory model, 720  
TLB (*see* Translation Lookaside Buffer)  
TLB miss, 461  
TN (*see* Twisted Nematic display)  
Token, 572  
Topology, 618

Touch screen, 113–115  
Towers of Hanoi, 417–420  
Towers of Hanoi for ARM, 418–420  
Towers of Hanoi for Core i7, 418–419  
Tracer, 693, 721–725  
Transaction layer, PCI Express, 227  
Transistor, invention, 19  
Transistor-transistor logic, 150  
Transition, 290  
Translation, 2  
Translation lookaside buffer, 331, 461  
Translator, 517  
Transmission control protocol, 576  
Transparent interrupt, 415  
Trap, 413, 706  
Trap handler, 413  
Tree, 620  
Tri-state device, 177  
TriMedia processor, 557–561  
Triple indirect block, 498  
True dependence, 297  
Truth table, 150  
TTL (*see* Transistor-Transistor Logic)  
Tuple, 642  
Tuple space, 642  
Twin, 642  
Twisted nematic display, 116  
Two's complement number, 675  
Two-pass assembler, 529–536  
Two-pass translator, 529  
TX-0, 19  
TX-2, 19

## U

U pipeline, 68  
UART (*see* Universal Asynchronous Receiver Transmitter)  
Ubiquitous computing, 28  
UCS (*see* Universal Character Set)  
UCS transformation format, 141  
UHCI (*see* Universal Host Controller Interface)  
UMA (*see* Uniform Memory Access computer)  
UMA (*see* Uniform Memory Access computer)  
Unicode, 138–141  
Unified cache, 84  
Uniform memory access computer, 592

Universal asynchronous receiver transmitter, 232  
 Universal character set, 141  
 Universal host controller interface, 231  
 Universal serial bus, 228–232  
   USB 2.0, 231  
   USB 3.0, 231  
 Universal synchronous asynchronous receiver transmitter, 232  
 UNIX, 482–485  
   virtual memory, 488–489  
 UNIX I/O, 492–498  
 UNIX process management, 503–506  
 Update strategy, cache, 600  
 USART (*see* Universal Synchronous Asynchronous Receiver Transmitter)  
 USB 2.0, 231  
 USB 3.0, 231  
 User mode, 347  
 UTF-8, 141–142

**V**

V pipeline, 68  
 Vacuum tube computers, 16–19  
 Vampire tap, 575  
 VAX, 14, 61  
 VCI (*see* Virtual Component Interconnect)  
 Vector processor, 71  
 Vector register, 71  
 Very large scale integration, 23  
 Very long instruction word, 555  
 VFP, 331  
 Video RAM, 117–118  
 Virtual 8086 mode, 352  
 Virtual address space, 440  
 Virtual circuit, 227  
 Virtual component interconnect, 574  
 Virtual cut through network, 625  
 Virtual machine, 3, 463–464  
 Virtual memory, 438–462  
   ARM, 460–462  
   compared to caching, 462  
   Core i7, 455–460  
   UNIX, 488–489  
   Windows 7, 490–492  
 Virtual organization, 652  
 Virtual register, 258

Virtual topology, 638  
 Virtual-machine control structure, 464  
 Virtualization, 464  
   hardware, 463–464  
 Virtuous circle, 29  
 VLIW (*see* Very Long Instruction Word)  
 VLSI (*see* Very Large Scale Integration)  
 VMCS (*see* Virtual-Machine Control Structure)  
 Volume table of contents, 105  
 Von Neumann, 18–19  
 Von Neumann machine, 18–19  
 VTOC (*see* Volume Table Of Contents)

**W**

Wait state, 192  
 WAN (*see* Wide Area Network)  
 WAR dependence, 318  
 Watson, Thomas, 19  
 WAW dependence, 318  
 Wax printer, 127  
 Weak consistency, 596–597  
 Wear leveling, 99  
 Weiser, Mark, 28  
 Weizac, 17  
 Whirlwind I, 14  
 Wide-area network, 575  
 Wiimote controller, 120–122  
 Wilkes, Maurice, 60  
 Win32 API, 487  
 Winchester disk, 88  
 Windows 7, 485–488  
 Windows 7 I/O, 498–503  
 Windows 7 process management, 506–509  
 Windows 7 virtual memory, 490–492  
 Windows drivers, 487  
 Windows new technology, 486  
 Windows NT (*see* Windows New Technology)  
 Wired-OR, 189  
 Word, 75  
 Word instruction, 701  
 Word register, 701  
 Working directory, 494  
 Working-set model, 443–446  
 Wozniak, Steve, 24  
 Write allocation, 310  
 Write back cache, 310

Write deferred cache, 310  
Write through cache, 310, 599  
Write-allocate policy, 600  
Write-back protocol, 601  
Write-once protocol, 601

## X

X Windows, 485  
x86, 25, 39–45, 347  
x86 architecture, 39  
XC2064, 14  
Xeon, 43

## Y

Y2K problem, 39, 359  
Yellow book, 101

## Z

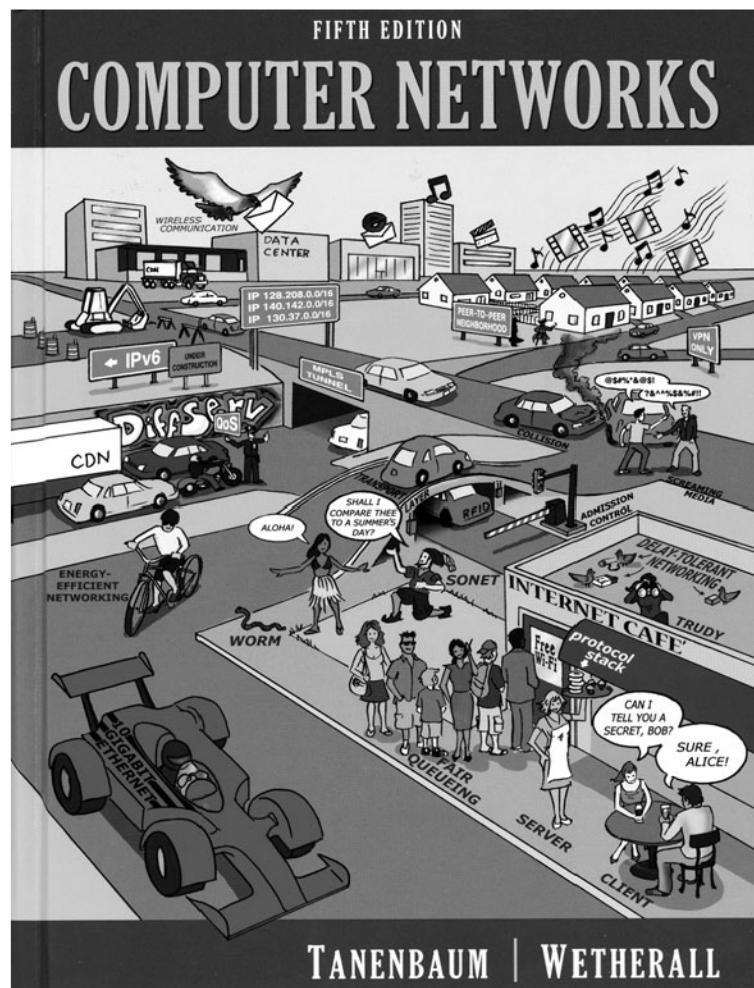
Z1, 14  
Zeroth-generation computers, 13–16  
Zilog Z8000, 61  
Zuse, Konrad, 15–16  
Zuse Z1, 14

*This page intentionally left blank*

**Also by Andrew S. Tanenbaum and David J. Wetherall**

Computer Networks, 5th ed.

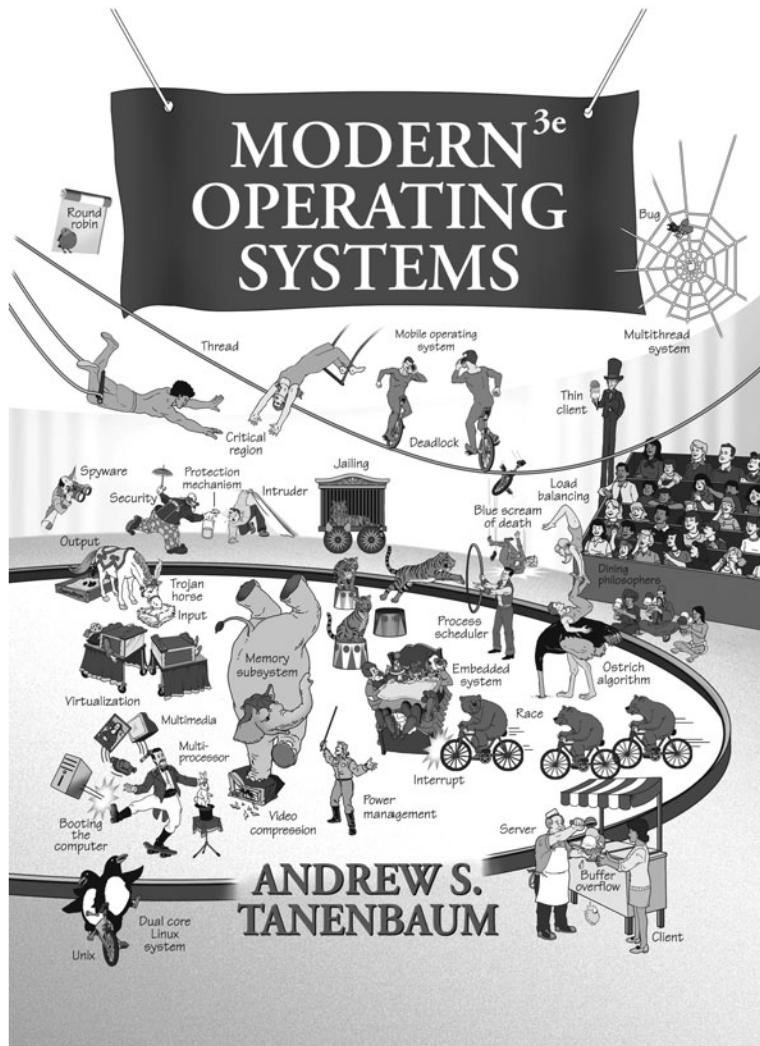
This widely read classic, with a fifth edition co-authored with David Wetherall, provides the ideal introduction to today's and tomorrow's networks. It explains in detail how modern networks are structured. Starting with the physical layer and working up to the application layer, the book covers a vast number of important topics, including wireless communication, fiber optics, data link protocols, Ethernet, routing algorithms, network performance, security, DNS, electronic mail, the World Wide Web, and multimedia. The book has especially thorough coverage of TCP/IP and the Internet.



Also by Andrew S. Tanenbaum

## Modern Operating Systems, 3rd ed.

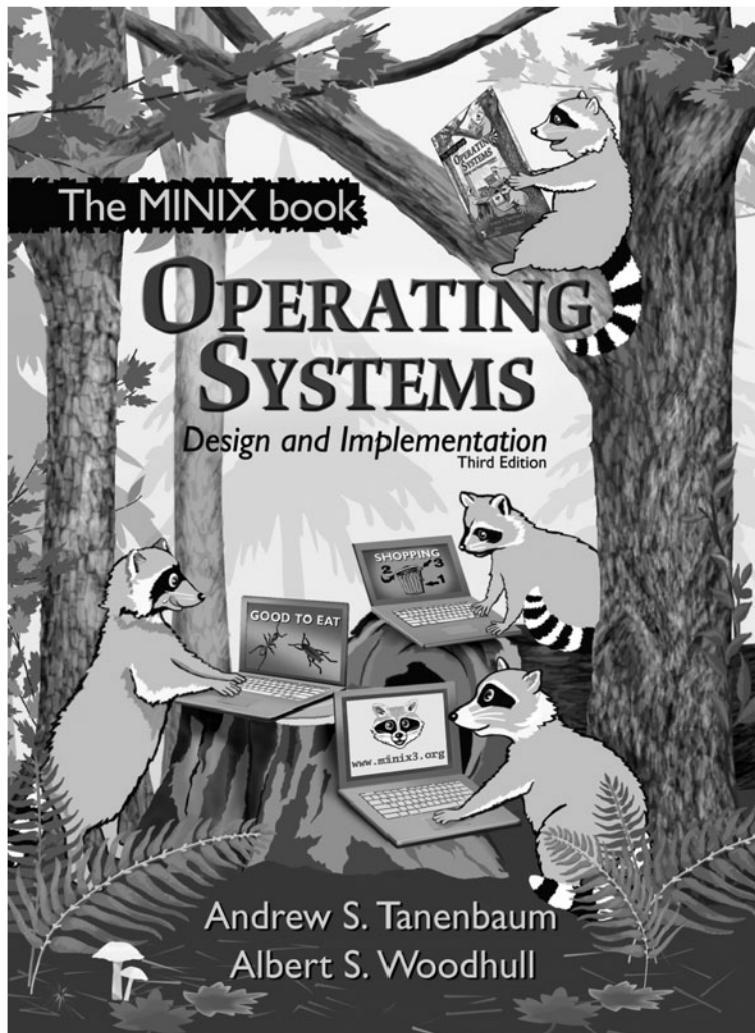
This worldwide best-seller incorporates the latest developments in operating systems. The book starts with chapters on the principles, including processes, memory management, file systems, I/O, and so on. Then it goes into three chapter-long case studies: Linux, Windows, and Symbian. Tanenbaum's experience as the designer of three operating systems (Amoeba, Globe, and MINIX) gives him a background few other authors can match, so the final chapter distills his long experience into advice for operating system designers.



Also by Andrew S. Tanenbaum and Albert S. Woodhull

## Operating Systems: Design and Implementation, 3rd ed.

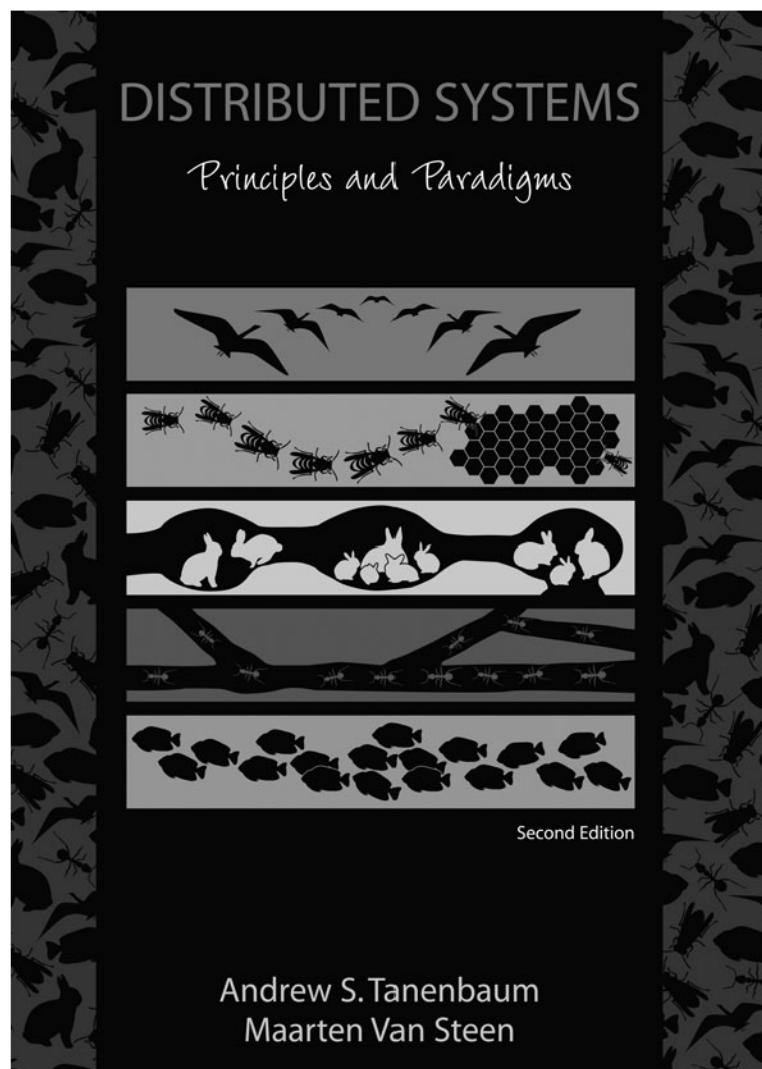
All other textbooks on operating systems are long on theory and short on practice. This one is different. In addition to the usual material on processes, memory management, file systems, I/O, and so on, it contains a CD-ROM with the source code (in C) of a small, but complete, POSIX-conformant operating system called MINIX 3 (see [www.minix3.org](http://www.minix3.org)). All the principles are illustrated by showing how they apply to MINIX 3. The reader can also compile, test, and experiment with MINIX 3, leading to in-depth knowledge of how an operating system really works.



Also by Andrew S. Tanenbaum and Maarten van Steen

## Distributed Systems: Principles and Paradigms, 2nd ed.

Distributed systems are becoming ever-more important in the world and this book explains their principles and illustrates them with numerous examples. Among the topics covered are architectures, processes, communication, naming, synchronization, consistency, fault tolerance, and security. Examples are taken from distributed object-based, file, Web-based, and coordination-based systems.



## CREDITS

p 8, "Hardware is just petrified software." © 2012 Karen Panetta. Reprinted by permission.

p 29, "Software is a gas. It expands to fill the container holding it.". "The Next Fifty Years of Software", talk prepared for ACM 97 by Nathan Myhrvold, 4/7/97. © 1997 Nathan Myhrvold. Reprinted by permission.

Figures 3-47, 3-48 © 2012 Texas Instruments. Reprinted courtesy of Texas Instruments.

Figures 3-49, Fig 3-50, 3-52, 4-50, 5-15 © 2012 Atmel Corporation. Reprinted by permission.

Figure 5-14 © ARM Ltd. Reprinted by permission.

## ABOUT THE AUTHORS

**Andrew S. Tanenbaum** has an S.B. degree from M.I.T. and a Ph.D. from the University of California at Berkeley. He is currently a Professor of Computer Science at the Vrije Universiteit where he has taught operating systems, networks, and related topics for over 30 years. His current research is on highly reliable operating systems although he has worked on compilers, distributed systems, security, and other topics over the years. These research projects have led to over 150 refereed papers in journals and conferences.

Prof. Tanenbaum has also (co)authored five books which have now appeared in 19 editions. The books have been translated into 21 languages, ranging from Basque to Thai and are used at universities all over the world. In all, there are 159 versions (language/edition combinations). See at [www.cs.vu.nl/~ast/book\\_covers](http://www.cs.vu.nl/~ast/book_covers) for images of them.

Prof. Tanenbaum has also produced a considerable volume of software, including the Amsterdam Compiler Kit (a retargetable portable compiler), Amoeba (an early distributed system used on LANs), and Globe (a wide-area distributed system).

He is also the author of MINIX, a small UNIX clone initially intended for use in student programming labs. It was the direct inspiration for Linux and the platform on which Linux was initially developed. The current version of MINIX, called MINIX 3, is now focused on being an extremely reliable and secure operating system. Prof. Tanenbaum will consider his work done when no computer is equipped with a reset button and no living person has ever experienced a system crash. MINIX 3 is an on-going open-source project to which you are invited to contribute. Go to [www.minix3.org](http://www.minix3.org) to download a free copy.

Tanenbaum is a Fellow of the ACM, a Fellow of the IEEE, and a member of the Royal Netherlands Academy of Arts and Sciences. He has also won numerous scientific prizes, including:

- 2010 TAA McGuffey Award for Computer Science and Engineering books
- 2007 IEEE James H. Mulligan, Jr. Education Medal
- 2002 TAA Texty Award for Computer Science and Engineering books
- 1997 SIGCSE Award for Outstanding Contributions to Computer Science Education
- 1994 ACM Karl V. Karlstrom Outstanding Educator Award

He also holds two honorary doctorates.

His home page on the World Wide Web can be found at <http://www.cs.vu.nl/~ast/>.

**Todd Austin** is a Professor of Electrical Engineering and Computer Science at the University of Michigan in Ann Arbor. His research interests include computer architecture, reliable and secure system design, hardware and software verification, and performance analysis tools and techniques. Prior to joining academia, Prof. Austin was a Senior Computer Architect in Intel's Microcomputer Research Labs, a product-oriented research laboratory in Hillsboro, Oregon. Prof. Austin is the first to take credit (but the last to accept blame) for creating the SimpleScalar Tool Set, a popular collection of computer architecture performance analysis tools. In 2002, Prof. Austin was a Sloan Research Fellow, and in 2007 he received the ACM Maurice Wilkes Award for “innovative contributions in Computer Architecture including the SimpleScalar Toolkit and the DIVA and Razor architectures.” Austin received his Ph.D. in Computer Science from the University of Wisconsin in 1996.